
Lab 01: A Gentle Introduction to Hadoop

CSC14118 Introduction to Big Data 20KHMT1

SmallData

2023-02-17

Contents

1	Lab 01: A Gentle Introduction to Hadoop	1
1.1	List of team members	1
1.2	Team's result	1
1.3	Team reflection	1
1.4	Setting up Single-node Hadoop Cluster	2
1.4.1	Step 1: Download java	2
1.4.2	Step 2: Create User for Hadoop and install openSSH	2
1.4.3	Step 3: Install Hadoop on Ubuntu	5
1.4.4	Step 4: Configuring Hadoop	8
1.4.5	Step 5: Start Hadoop Cluster	13
1.5	Introduction to MapReduce	14
1.6	Running a warm-up problem: Word Count	17
1.6.1	Step 0: Install Eclipse on Ubuntu (if you had installed, please go to next step)	18
1.6.2	Step 1: Create new Java project	18
1.6.3	Step 2: Delete file <i>module-info.java</i>	18
1.6.4	Step 3: Create Java package	21
1.6.5	Step 4: Create Java class	21
1.6.6	Step 5: Paste WordCount code to the <i>WordCount.java</i> file just created	21
1.6.7	Step 6: Configure build path for the project	24
1.6.8	Step 7: Export to JAR file	24
1.6.9	Step 8: Prepare to run MapReduce	28
1.6.10	Step 9: Run MapReduce	31
1.7	Bonus	31
1.7.1	4.1 Extended Word Count: Unhealthy relationships	31
1.7.2	4.2 Setting up Fully Distributed Mode	34
1.8	References	34

1 Lab 01: A Gentle Introduction to Hadoop

1.1 List of team members

ID	Full Name
20120366	Pham Phu Hoang Son
20120391	Ha Xuan Truong
20120393	Huynh Minh Tu
20120468	Nguyen Van Hai

1.2 Team's result

Section	Complete
Setting up SNC	100%
Introduction to MapReduce	100%
Running a warm-up problem: Word Count	100%
Bonus - Extended Word Count: Unhealthy relationships	100%
Bonus - Setting up Fully Distributed Mode	0%

1.3 Team reflection

Does your journey to the deadline have any bugs? How have you overcome it?

During the journey towards the deadline, we encountered several bugs that were related to Ubuntu, Hadoop installation, and errors while running Hadoop MapReduce jobs. In order to overcome these

challenges, we had to invest more time and effort. We also conducted research by reading documentation and watching tutorial videos. These resources provided us with useful insights and ideas for troubleshooting the issues. Additionally, we scheduled some online meetings to discuss and solve the problems together. Through these efforts, we were able to solve most of the problems we encountered and successfully complete the project.

What have you learned after this process?

Firstly, we learned the importance of clear communication among team members to ensure that everyone is on the same page and that tasks are completed efficiently. We also learned the importance of testing and debugging to ensure that any errors are caught and resolved early on in the process.

Secondly, we learned the importance of time management and task prioritization, as we encountered some unexpected challenges during the installation and setup process. This made it necessary for us to adjust our timeline and focus on the most critical tasks first.

Lastly, we learned the importance of continuous learning and self-improvement. We encountered some roadblocks that required us to do additional research and seek out new solutions, which allowed us to expand our knowledge and skills in Hadoop and MapReduce.

1.4 Setting up Single-node Hadoop Cluster

1.4.1 Step 1: Download java

1. The default Ubuntu repositories contain Java 8 and Java 11 both. Use the following command to install it.

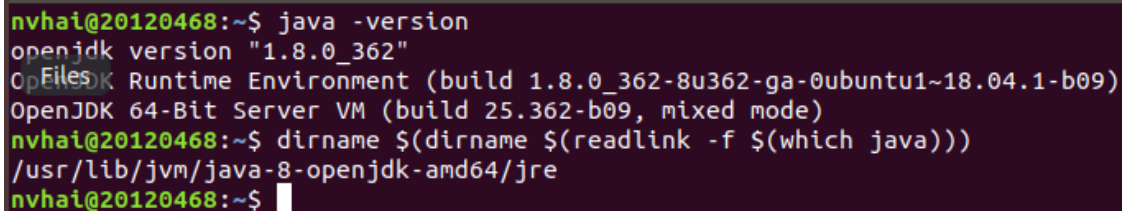
```
sudo apt update && sudo apt install openjdk-8-jdk
```

2. Once you have successfully installed it, check the current Java version:

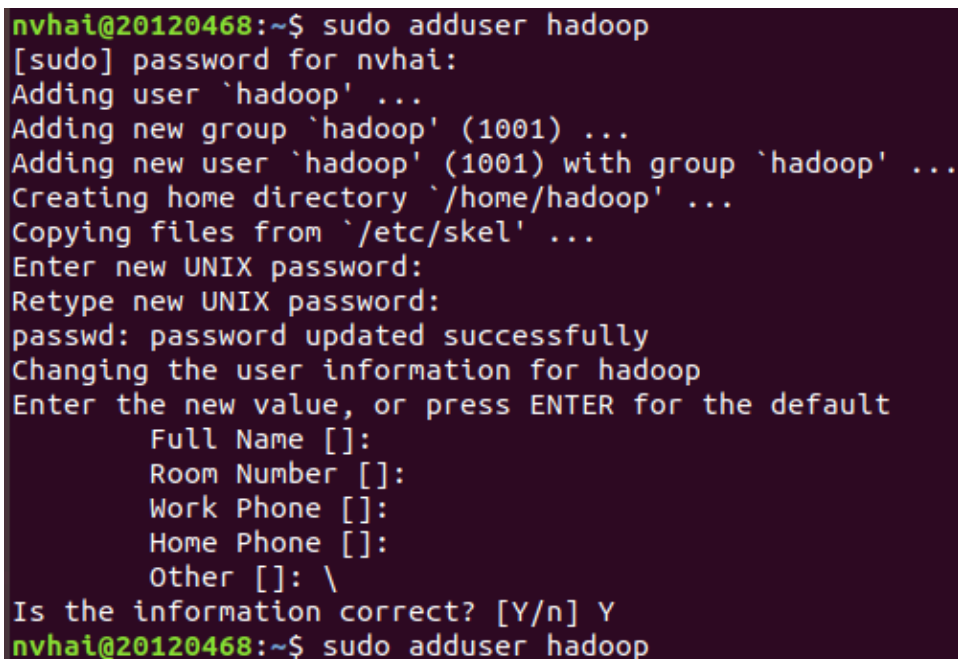
```
java -version
```

1.4.2 Step 2: Create User for Hadoop and install openSSH

1. Run the following command to create a new user with the name “hadoop”:

A terminal window with a dark purple background. The prompt is 'nvhai@20120468:~\$'. The user enters 'java -version'. The output is 'openjdk version "1.8.0_362"', 'OpenJDK Runtime Environment (build 1.8.0_362-8u362-ga-0ubuntu1~18.04.1-b09)', and 'OpenJDK 64-Bit Server VM (build 25.362-b09, mixed mode)'. The user then enters 'dirname \$(dirname \$(readlink -f \$(which java)))'. The output is '/usr/lib/jvm/java-8-openjdk-amd64/jre'. The prompt returns to 'nvhai@20120468:~\$'.

```
nvhai@20120468:~$ java -version
openjdk version "1.8.0_362"
OpenJDK Runtime Environment (build 1.8.0_362-8u362-ga-0ubuntu1~18.04.1-b09)
OpenJDK 64-Bit Server VM (build 25.362-b09, mixed mode)
nvhai@20120468:~$ dirname $(dirname $(readlink -f $(which java)))
/usr/lib/jvm/java-8-openjdk-amd64/jre
nvhai@20120468:~$
```

Figure 1.1: Download javaA terminal window with a dark purple background. The prompt is 'nvhai@20120468:~\$'. The user enters 'sudo adduser hadoop'. The output shows the process of adding a user: '[sudo] password for nvhai:', 'Adding user `hadoop` ...', 'Adding new group `hadoop` (1001) ...', 'Adding new user `hadoop` (1001) with group `hadoop` ...', 'Creating home directory `/home/hadoop` ...', 'Copying files from `/etc/skel` ...', 'Enter new UNIX password:', 'Retype new UNIX password:', 'passwd: password updated successfully', 'Changing the user information for hadoop', 'Enter the new value, or press ENTER for the default', 'Full Name []:', 'Room Number []:', 'Work Phone []:', 'Home Phone []:', 'Other []: \', 'Is the information correct? [Y/n] Y'. The prompt returns to 'nvhai@20120468:~\$'.

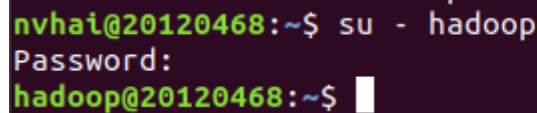
```
nvhai@20120468:~$ sudo adduser hadoop
[sudo] password for nvhai:
Adding user `hadoop` ...
Adding new group `hadoop` (1001) ...
Adding new user `hadoop` (1001) with group `hadoop` ...
Creating home directory `/home/hadoop` ...
Copying files from `/etc/skel` ...
Enter new UNIX password:
Retype new UNIX password:
passwd: password updated successfully
Changing the user information for hadoop
Enter the new value, or press ENTER for the default
  Full Name []:
   Room Number []:
    Work Phone []:
   Home Phone []:
        Other []: \
Is the information correct? [Y/n] Y
nvhai@20120468:~$ sudo adduser hadoop
```

Figure 1.2: Create new user

```
sudo adduser hadoop
```

2. Switch to the newly created hadoop user:

```
su - hadoop
```

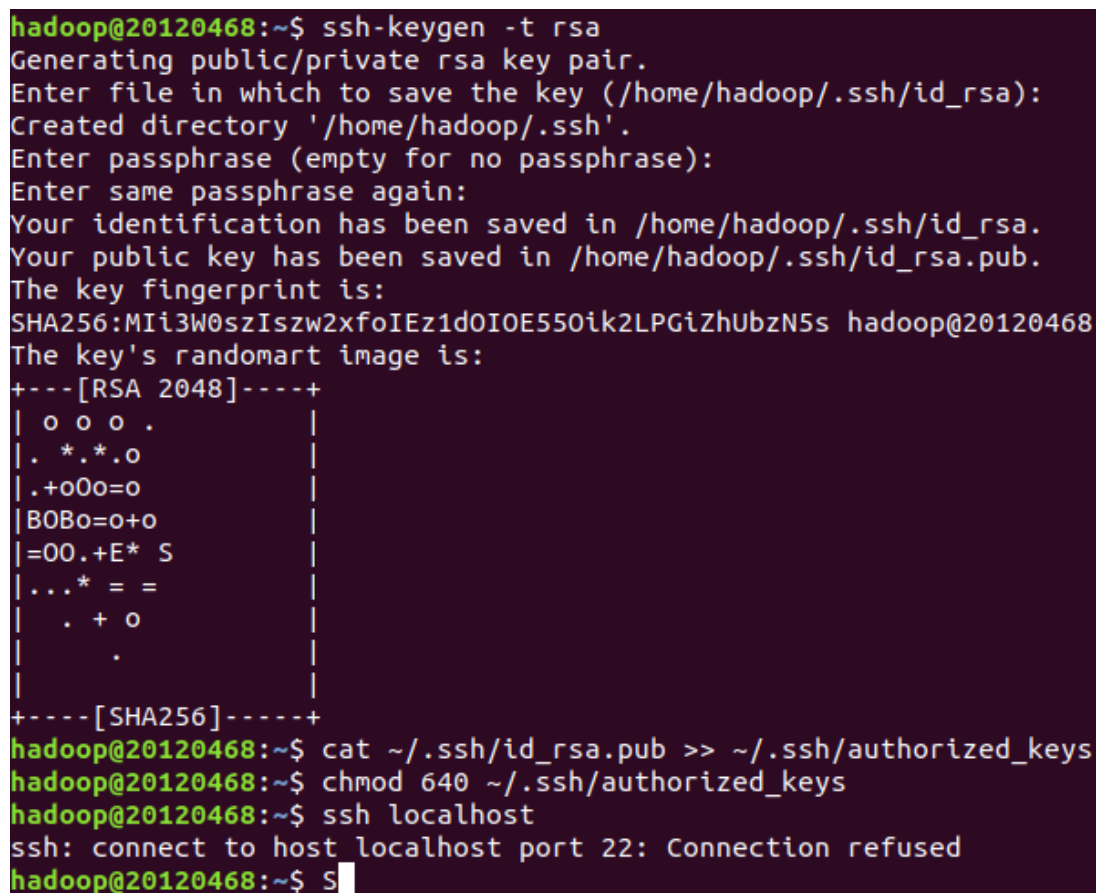


```
nvhai@20120468:~$ su - hadoop
Password:
hadoop@20120468:~$
```

Figure 1.3: Change to hadoop user

3. Now configure password-less SSH access for the newly created hadoop user. Generate an SSH keypair first:

```
ssh-keygen -t rsa
```



```
hadoop@20120468:~$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/home/hadoop/.ssh/id_rsa):
Created directory '/home/hadoop/.ssh'.
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/hadoop/.ssh/id_rsa.
Your public key has been saved in /home/hadoop/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:MIi3W0szIszw2xfoIEz1d0IOE550ik2LPGiZhUbzN5s hadoop@20120468
The key's randomart image is:
+---[RSA 2048]---+
| o o o .          |
| . *. *. o        |
| .+oOo=o         |
| B0Bo=o+o        |
| =00.+E* S       |
| ...* = =        |
| . + o           |
| .               |
+-----[SHA256]-----+
hadoop@20120468:~$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
hadoop@20120468:~$ chmod 640 ~/.ssh/authorized_keys
hadoop@20120468:~$ ssh localhost
ssh: connect to host localhost port 22: Connection refused
hadoop@20120468:~$
```

Figure 1.4: OpenSSH

4. Copy the generated public key to the authorized key file and set the proper permissions:

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
chmod 640 ~/.ssh/authorized_keys
```

5. Now try to SSH to the localhost

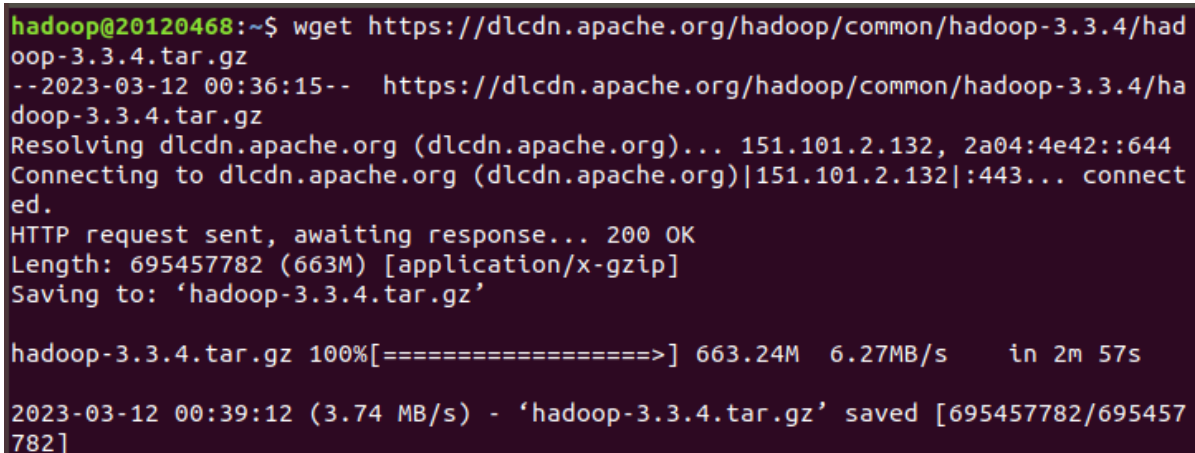
```
ssh localhost
```

You will be asked to authenticate hosts by adding RSA keys to known hosts. Type yes and hit Enter to authenticate the localhost.

1.4.3 Step 3: Install Hadoop on Ubuntu

1. Use the following command to download Hadoop 3.3.4

```
wget https://dlcdn.apache.org/hadoop/common/hadoop-3.3.4/hadoop-3.3.4.tar.gz
```



```
hadoop@20120468:~$ wget https://dlcdn.apache.org/hadoop/common/hadoop-3.3.4/hadoop-3.3.4.tar.gz
--2023-03-12 00:36:15-- https://dlcdn.apache.org/hadoop/common/hadoop-3.3.4/hadoop-3.3.4.tar.gz
Resolving dlcdn.apache.org (dlcdn.apache.org)... 151.101.2.132, 2a04:4e42::644
Connecting to dlcdn.apache.org (dlcdn.apache.org)|151.101.2.132|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 695457782 (663M) [application/x-gzip]
Saving to: 'hadoop-3.3.4.tar.gz'

hadoop-3.3.4.tar.gz 100%[=====>] 663.24M 6.27MB/s in 2m 57s

2023-03-12 00:39:12 (3.74 MB/s) - 'hadoop-3.3.4.tar.gz' saved [695457782/695457782]
```

Figure 1.5: download hadoop

2. Once you've downloaded the file, you can unzip it to a folder on your hard drive

```
tar xzf hadoop-3.3.4.tar.gz
```

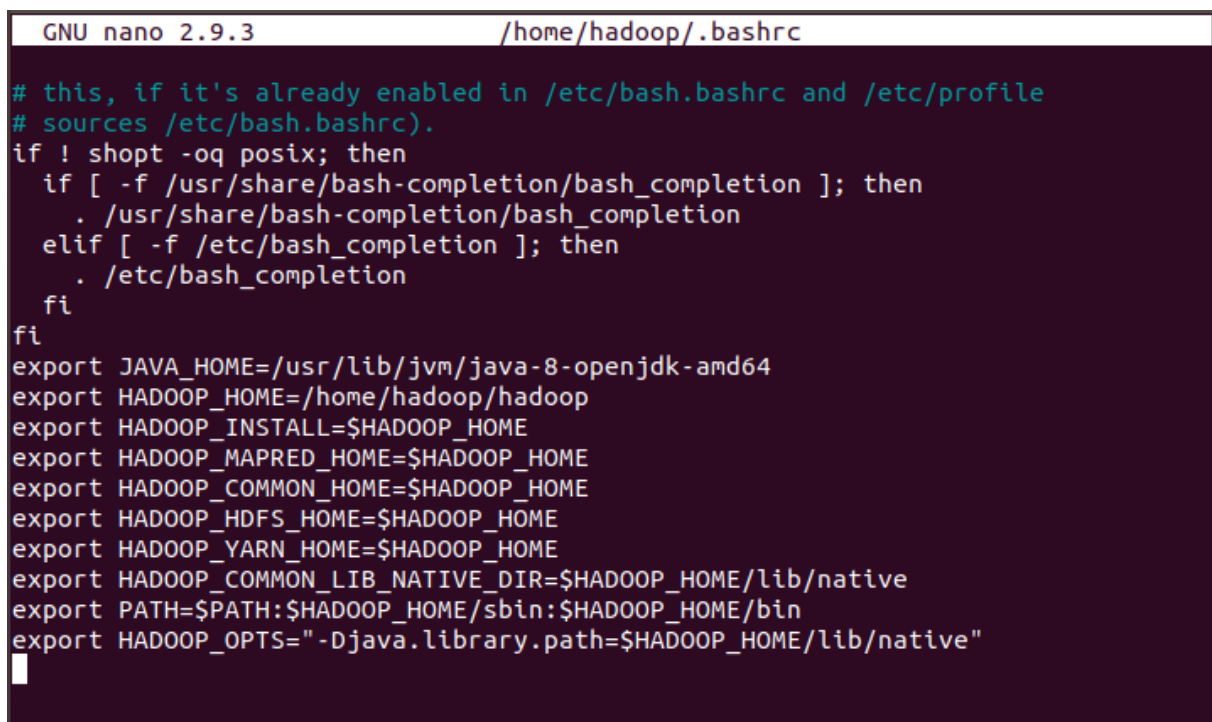
3. Rename the extracted folder to remove version information. This is an optional step, but if you don't want to rename, then adjust the remaining configuration paths.

```
mv hadoop-3.3.4 hadoop
```

- Next, you will need to configure Hadoop and Java Environment Variables on your system. Open the ~/.bashrc file in your favorite text editor:

```
nano ~/.bashrc
```

Append the below lines to the file. You can find the JAVA_HOME location by running `dirname $(dirname $(readlink -f $(which java)))` command on the terminal.



```
GNU nano 2.9.3 /home/hadoop/.bashrc

# this, if it's already enabled in /etc/bash.bashrc and /etc/profile
# sources /etc/bash.bashrc).
if ! shopt -oq posix; then
  if [ -f /usr/share/bash-completion/bash_completion ]; then
    . /usr/share/bash-completion/bash_completion
  elif [ -f /etc/bash_completion ]; then
    . /etc/bash_completion
  fi
fi
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
export HADOOP_HOME=/home/hadoop/hadoop
export HADOOP_INSTALL=$HADOOP_HOME
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_COMMON_HOME=$HADOOP_HOME
export HADOOP_HDFS_HOME=$HADOOP_HOME
export HADOOP_YARN_HOME=$HADOOP_HOME
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native
export PATH=$PATH:$HADOOP_HOME/sbin:$HADOOP_HOME/bin
export HADOOP_OPTS="-Djava.library.path=$HADOOP_HOME/lib/native"
```

Figure 1.6: setup-environment

Save the file and close it.

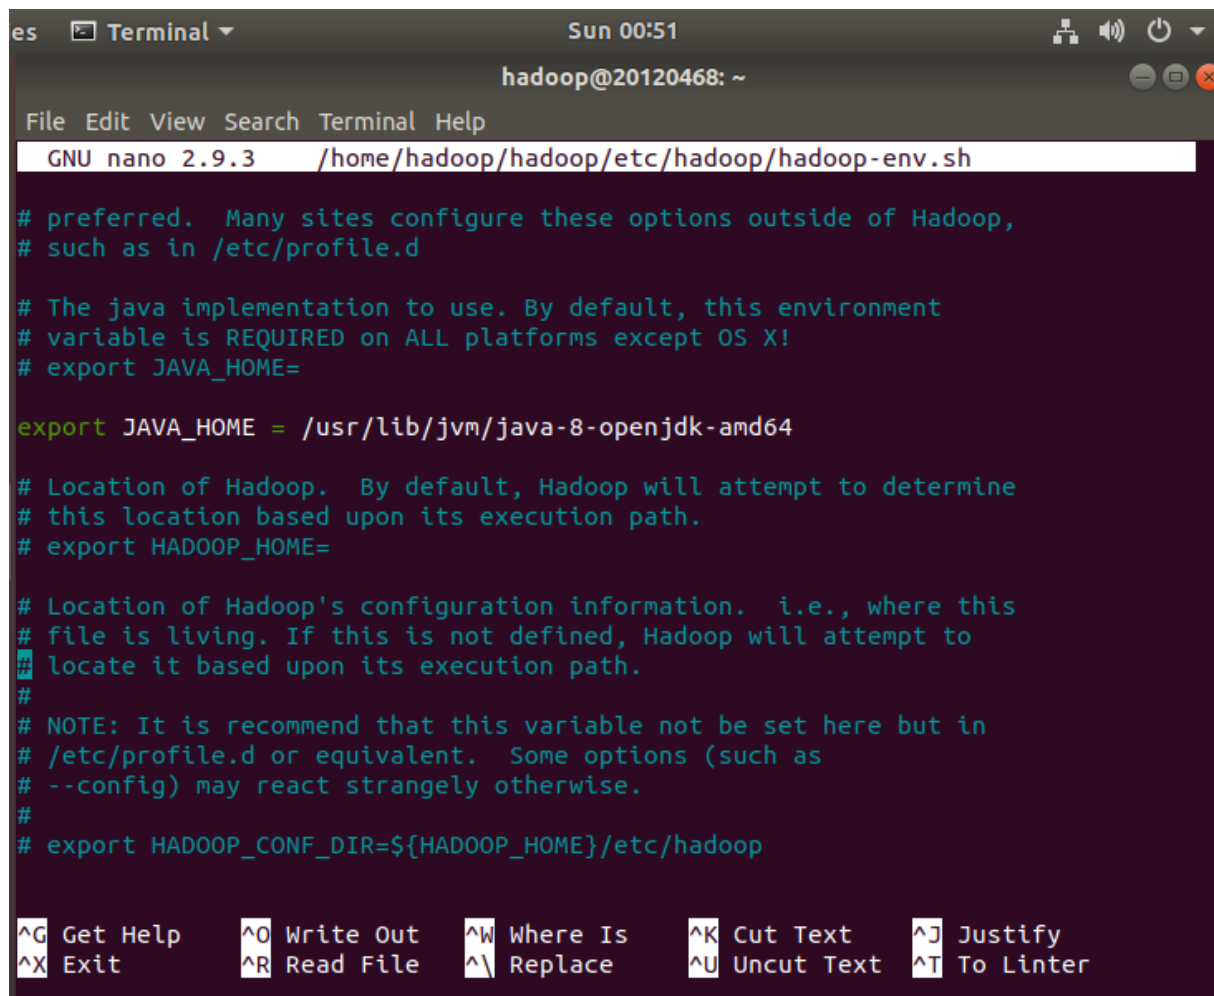
- Load the above configuration in the current environment

```
source ~/.bashrc
```

- You also need to configure JAVA_HOME in hadoop-env.sh file. Edit the Hadoop environment variable file in the text editor:

```
nano $HADOOP_HOME/etc/hadoop/hadoop-env.sh
```

Search for the “export JAVA_HOME” and configure it with the value found in step 1. See the below screenshot:



```
es  Terminal ▾  Sun 00:51  hadoop@20120468: ~
File Edit View Search Terminal Help
GNU nano 2.9.3 /home/hadoop/hadoop/etc/hadoop/hadoop-env.sh

# preferred.  Many sites configure these options outside of Hadoop,
# such as in /etc/profile.d

# The java implementation to use.  By default, this environment
# variable is REQUIRED on ALL platforms except OS X!
# export JAVA_HOME=

export JAVA_HOME = /usr/lib/jvm/java-8-openjdk-amd64

# Location of Hadoop.  By default, Hadoop will attempt to determine
# this location based upon its execution path.
# export HADOOP_HOME=

# Location of Hadoop's configuration information.  i.e., where this
# file is living.  If this is not defined, Hadoop will attempt to
# locate it based upon its execution path.
#
# NOTE: It is recommend that this variable not be set here but in
# /etc/profile.d or equivalent.  Some options (such as
# --config) may react strangely otherwise.
#
# export HADOOP_CONF_DIR=${HADOOP_HOME}/etc/hadoop

^G Get Help  ^O Write Out  ^W Where Is  ^K Cut Text  ^J Justify
^X Exit      ^R Read File  ^\ Replace   ^U Uncut Text ^T To Linter
```

Figure 1.7: setup-hadoop-env

Save the file and close it.

1.4.4 Step 4: Configuring Hadoop

Next is to configure Hadoop configuration files available under etc directory.

1. First, you will need to create the namenode and datanode directories inside the Hadoop user home directory. Run the following command to create both directories:

```
mkdir -p ~/hadoopdata/hdfs/{namenode,datanode}
```

2. Next, edit the core-site.xml file and update with your system hostname:

```
nano $HADOOP_HOME/etc/hadoop/core-site.xml
```

Change the following name as per your system hostname:

Save and close the file.

3. Then, edit the hdfs-site.xml file

```
nano $HADOOP_HOME/etc/hadoop/core-site.xml
```

Change the NameNode and DataNode directory paths as shown below:

Save and close the file.

4. Then, edit the mapred-site.xml file

```
nano $HADOOP_HOME/etc/hadoop/mapred-site.xml
```

Make the following changes:

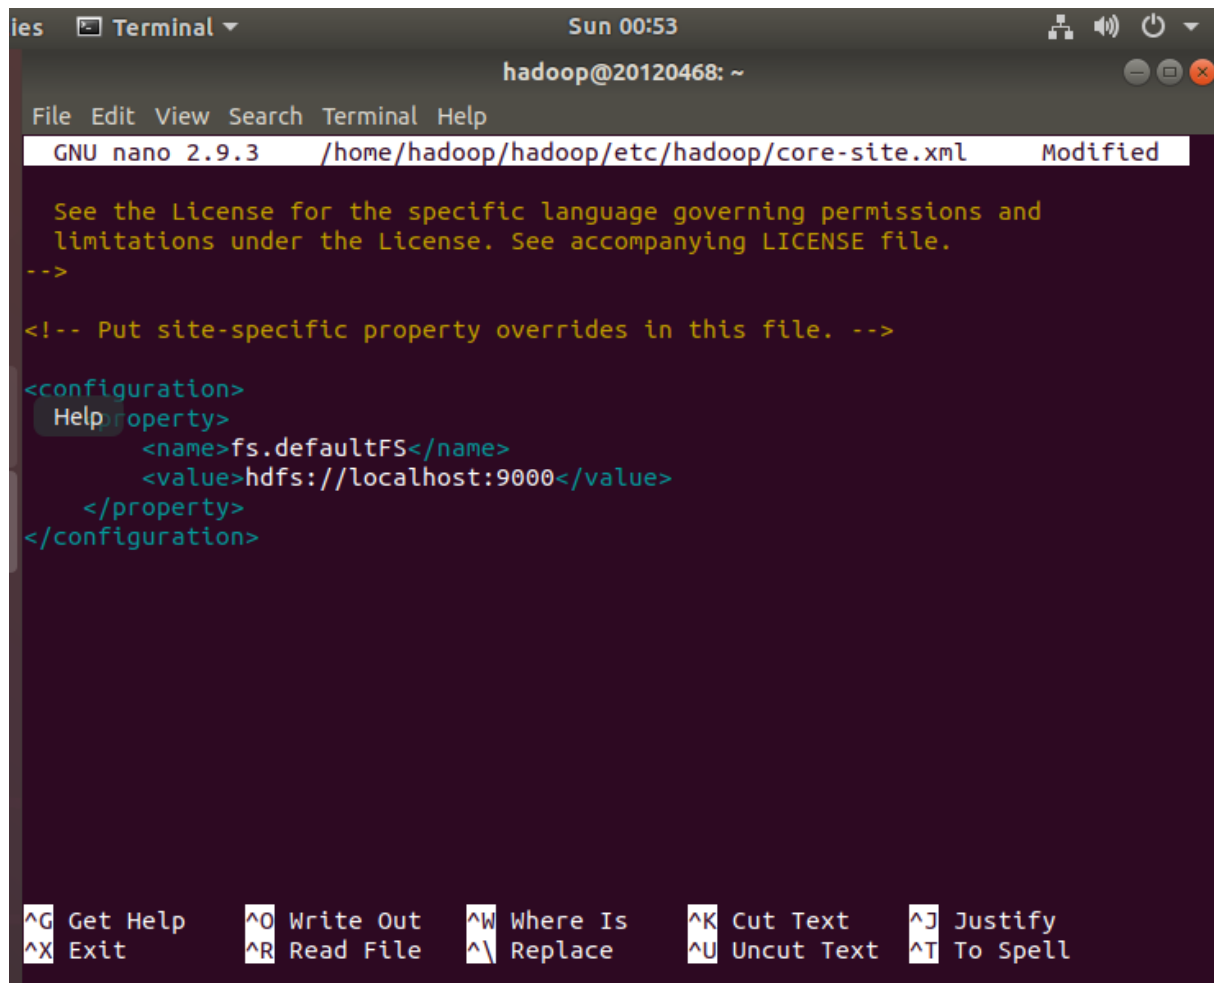
Save and close the file.

5. Then, edit the yarn-site.xml file

```
nano $HADOOP_HOME/etc/hadoop/yarn-site.xml
```

Make the following changes:

Save and close the file.



```
ies Terminal Sun 00:53
hadoop@20120468: ~
File Edit View Search Terminal Help
GNU nano 2.9.3 /home/hadoop/hadoop/etc/hadoop/core-site.xml Modified

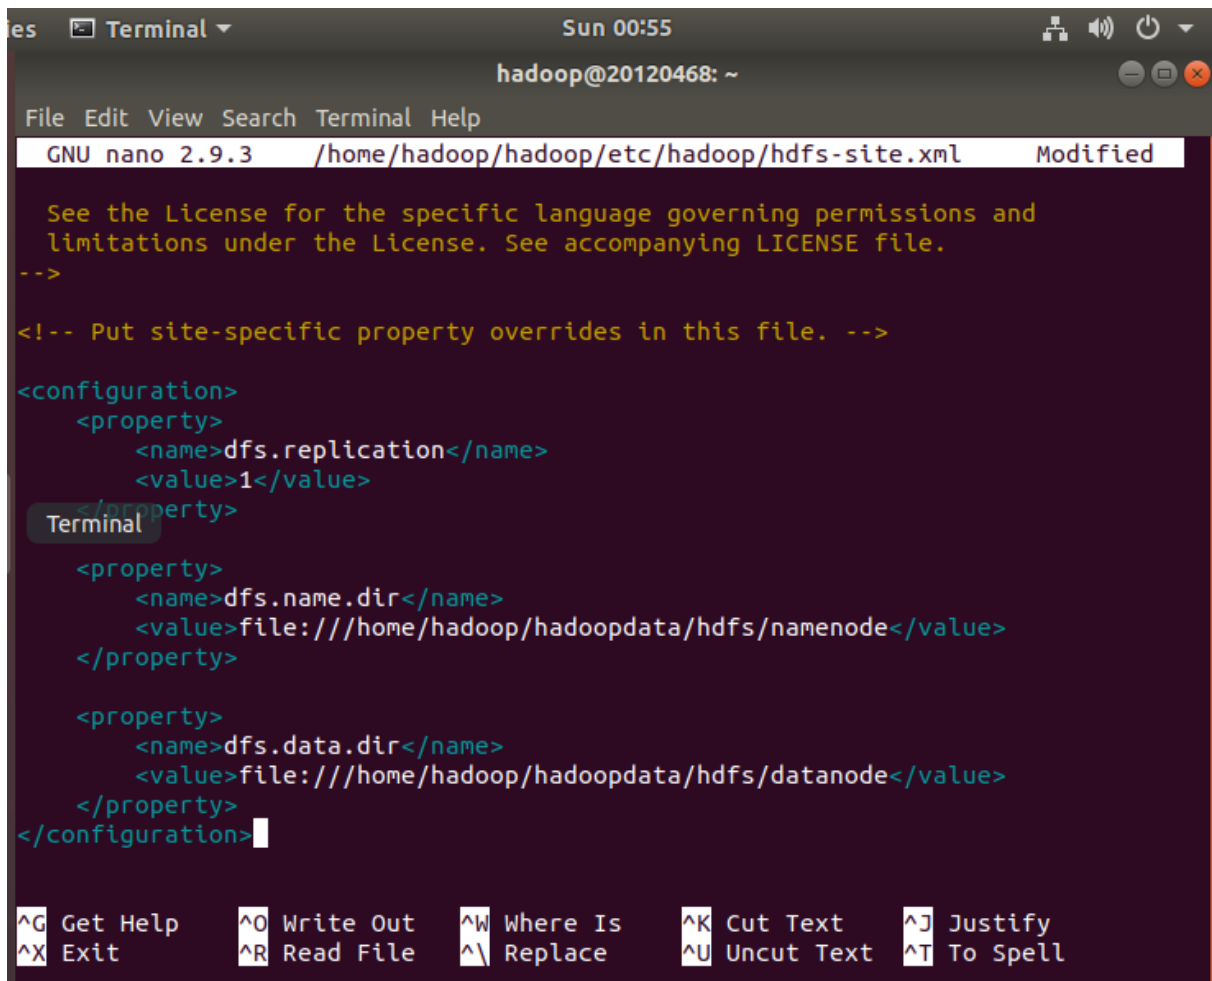
See the License for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>

^G Get Help    ^O Write Out  ^W Where Is   ^K Cut Text   ^J Justify
^X Exit        ^R Read File  ^\ Replace    ^U Uncut Text ^T To Spell
```

Figure 1.8: setup-core-site



```
es Terminal Sun 00:55
hadoop@20120468: ~
File Edit View Search Terminal Help
GNU nano 2.9.3 /home/hadoop/hadoop/etc/hadoop/hdfs-site.xml Modified

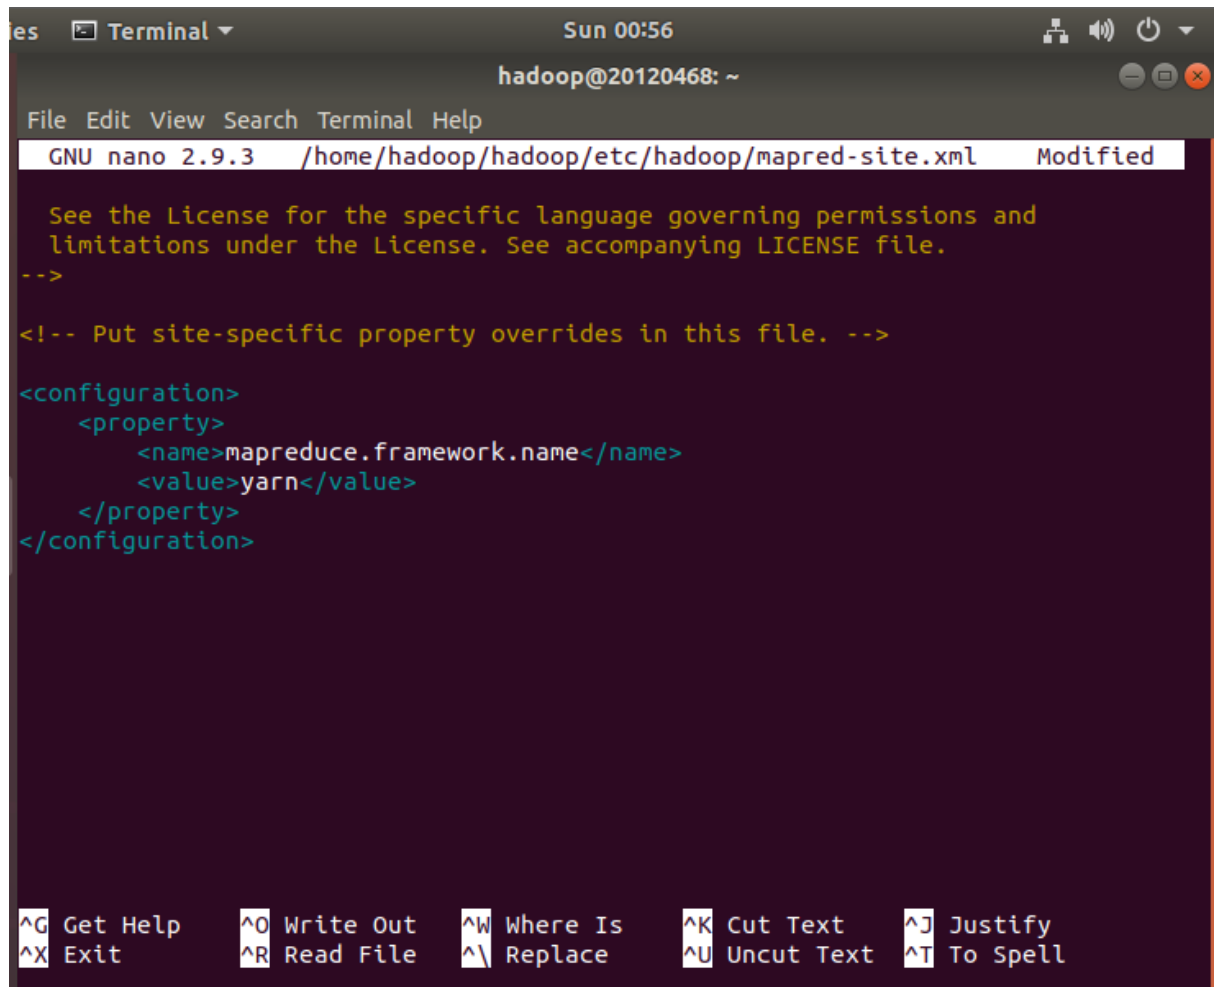
See the License for the specific language governing permissions and
limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
  <property>
    <name>dfs.name.dir</name>
    <value>file:///home/hadoop/hadoopdata/hdfs/namenode</value>
  </property>
  <property>
    <name>dfs.data.dir</name>
    <value>file:///home/hadoop/hadoopdata/hdfs/datanode</value>
  </property>
</configuration>

^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify
^X Exit ^R Read File ^\ Replace ^U Uncut Text ^T To Spell
```

Figure 1.9: setup-hdfs-site



```
es  Terminal ▾                               Sun 00:56
                                           hadoop@20120468: ~
File Edit View Search Terminal Help
GNU nano 2.9.3 /home/hadoop/hadoop/etc/hadoop/mapred-site.xml Modified

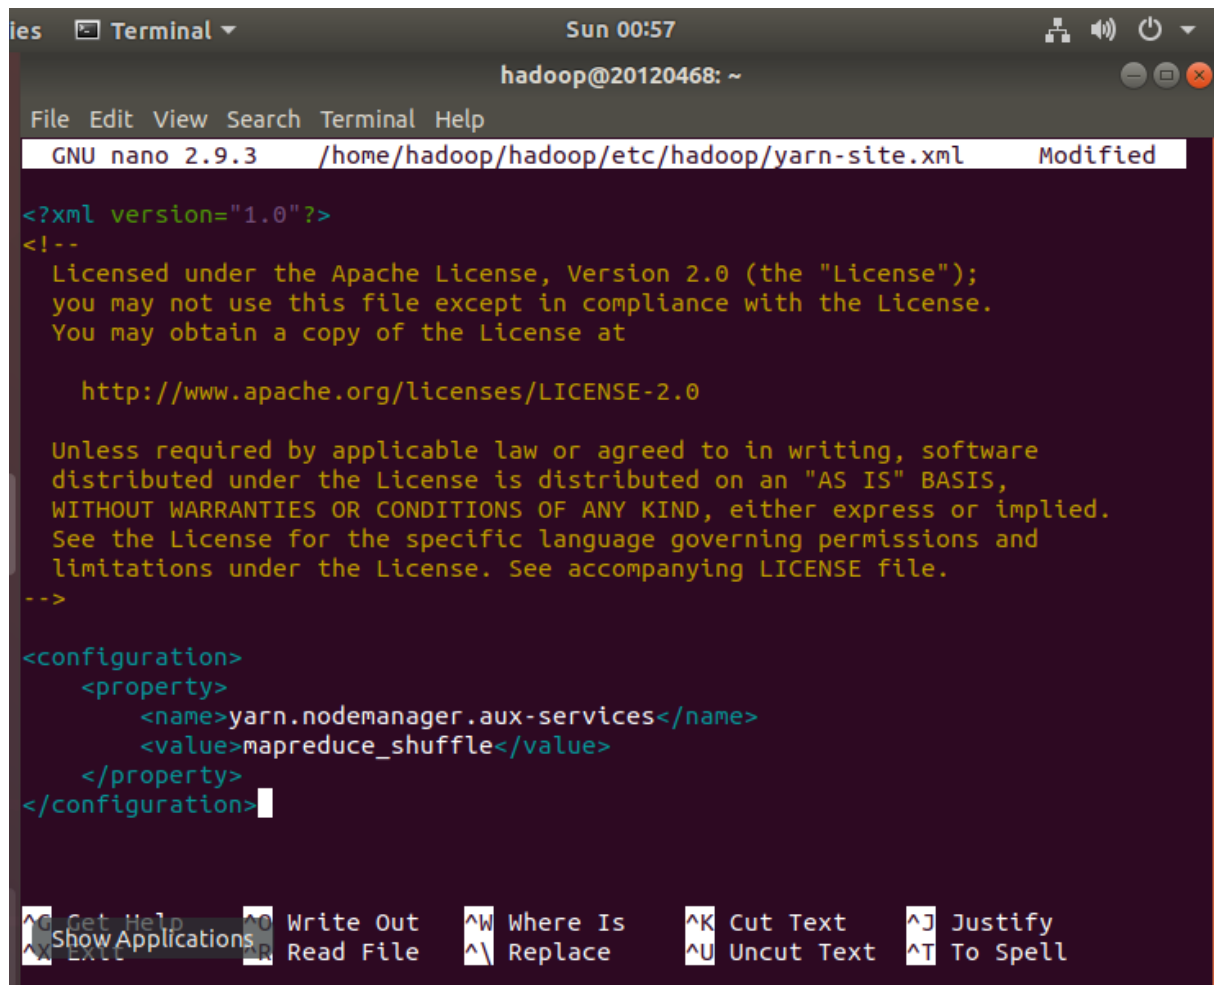
  See the License for the specific language governing permissions and
  limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
</configuration>

^G Get Help      ^O Write Out    ^W Where Is     ^K Cut Text     ^J Justify
^X Exit          ^R Read File    ^_ Replace      ^U Uncut Text   ^T To Spell
```

Figure 1.10: setup-mapred-site



```
ies Terminal ▾ Sun 00:57
hadoop@20120468: ~
File Edit View Search Terminal Help
GNU nano 2.9.3 /home/hadoop/hadoop/etc/hadoop/yarn-site.xml Modified

<?xml version="1.0"?>
<!--
  Licensed under the Apache License, Version 2.0 (the "License");
  you may not use this file except in compliance with the License.
  You may obtain a copy of the License at

      http://www.apache.org/licenses/LICENSE-2.0

  Unless required by applicable law or agreed to in writing, software
  distributed under the License is distributed on an "AS IS" BASIS,
  WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
  See the License for the specific language governing permissions and
  limitations under the License. See accompanying LICENSE file.
-->

<configuration>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
</configuration>
```

^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify
^X Show Applications ^R Read File ^\ Replace ^U Uncut Text ^T To Spell

Figure 1.11: setup-yarn-site

1.4.5 Step 5: Start Hadoop Cluster

Then start the Hadoop cluster with the following command

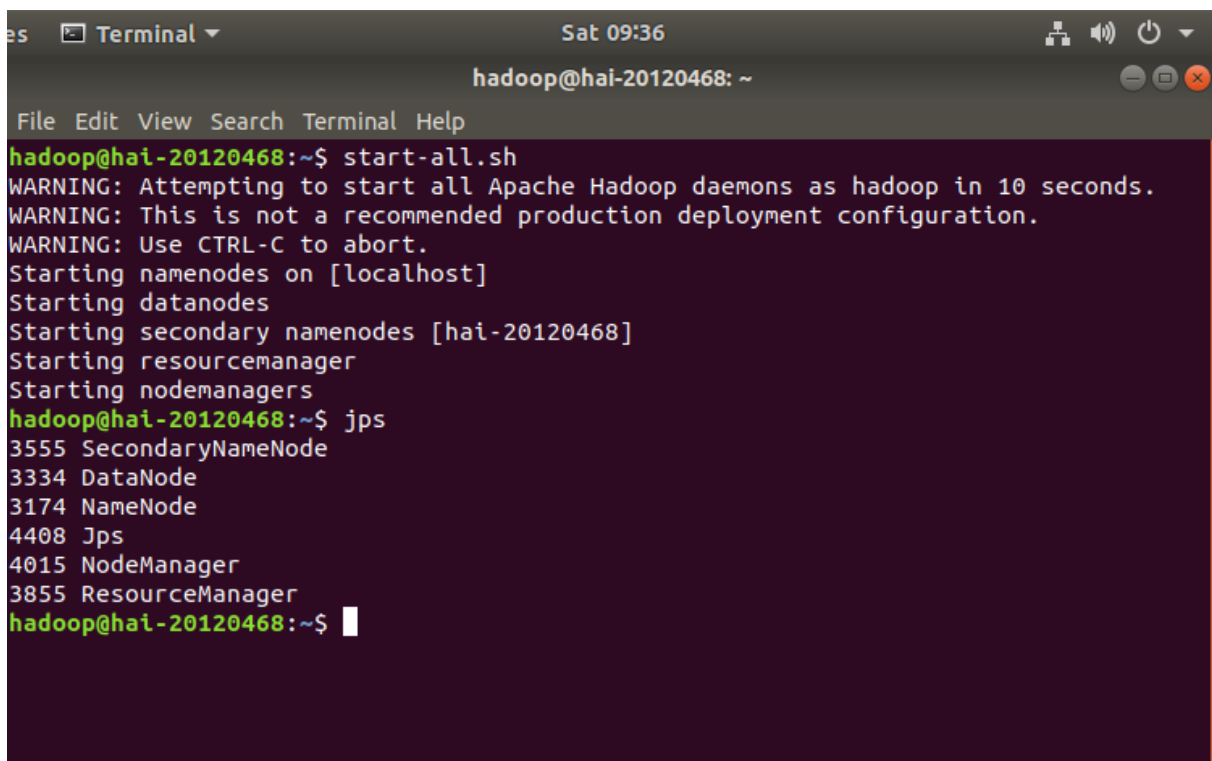
```
start-all.sh
```

Check jps

```
jps
```

Completed screenshots of the members:

20120468 - Nguyen Van Hai

A terminal window titled 'Terminal' with a dark background. The prompt is 'hadoop@hai-20120468: ~'. The user enters 'start-all.sh'. The output shows several warnings and then the start of various Hadoop daemons: namenodes on [localhost], datanodes, secondary namenodes [hai-20120468], resourcemanager, and nodemanagers. Then the user enters 'jps', and the output lists the running processes with their PIDs: 3555 SecondaryNameNode, 3334 DataNode, 3174 NameNode, 4408 Jps, 4015 NodeManager, and 3855 ResourceManager.

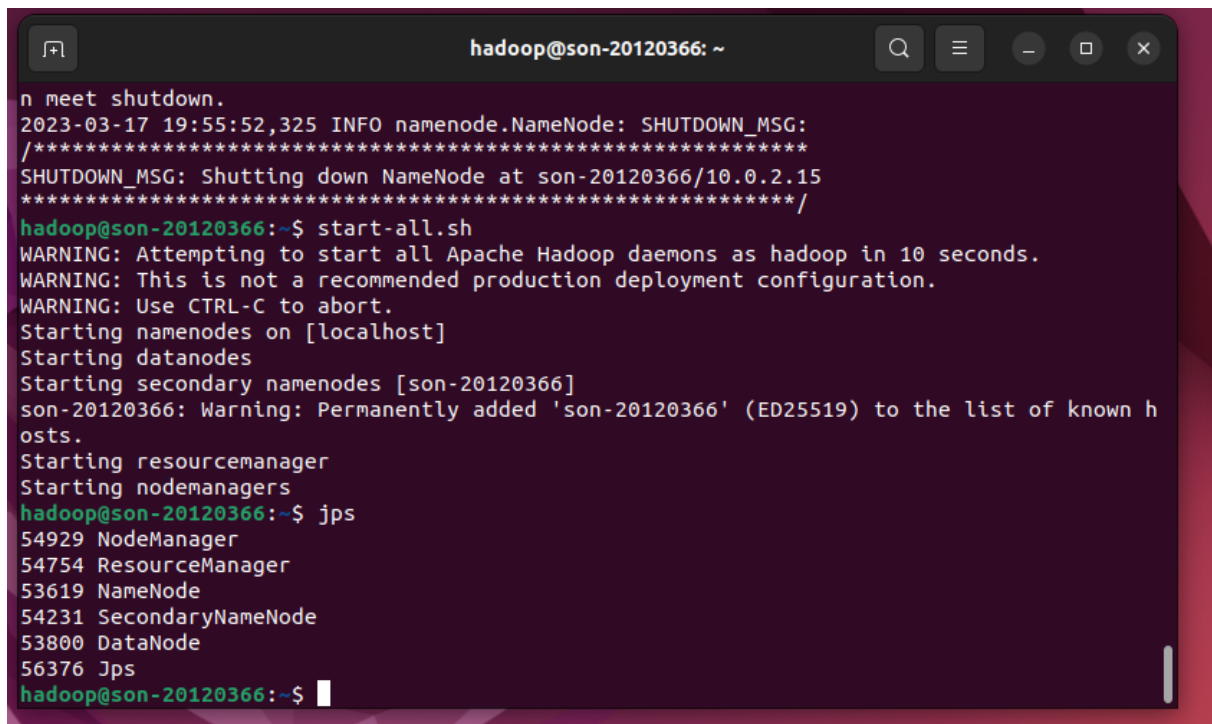
```
es Terminal Sat 09:36
hadoop@hai-20120468: ~
File Edit View Search Terminal Help
hadoop@hai-20120468:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [hai-20120468]
Starting resourcemanager
Starting nodemanagers
hadoop@hai-20120468:~$ jps
3555 SecondaryNameNode
3334 DataNode
3174 NameNode
4408 Jps
4015 NodeManager
3855 ResourceManager
hadoop@hai-20120468:~$
```

Figure 1.12: 20120468 done

20120366 - Pham Phu Hoang Son

20120391 - Ha Xuan Truong

20120393 - Huynh Minh Tu

A terminal window titled 'hadoop@son-20120366: ~' with standard window controls. The terminal shows a sequence of events: a shutdown message from a previous session, the execution of 'start-all.sh' which displays several warnings and status messages about starting Hadoop daemons, and the execution of 'jps' which lists the running processes: NodeManager, ResourceManager, NameNode, SecondaryNameNode, and DataNode.

```
hadoop@son-20120366: ~  
n meet shutdown.  
2023-03-17 19:55:52,325 INFO namenode.NameNode: SHUTDOWN_MSG:  
/*****  
SHUTDOWN_MSG: Shutting down NameNode at son-20120366/10.0.2.15  
*****/  
hadoop@son-20120366:~$ start-all.sh  
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.  
WARNING: This is not a recommended production deployment configuration.  
WARNING: Use CTRL-C to abort.  
Starting namenodes on [localhost]  
Starting datanodes  
Starting secondary namenodes [son-20120366]  
son-20120366: Warning: Permanently added 'son-20120366' (ED25519) to the list of known h  
osts.  
Starting resourcemanager  
Starting nodemanagers  
hadoop@son-20120366:~$ jps  
54929 NodeManager  
54754 ResourceManager  
53619 NameNode  
54231 SecondaryNameNode  
53800 DataNode  
56376 Jps  
hadoop@son-20120366:~$
```

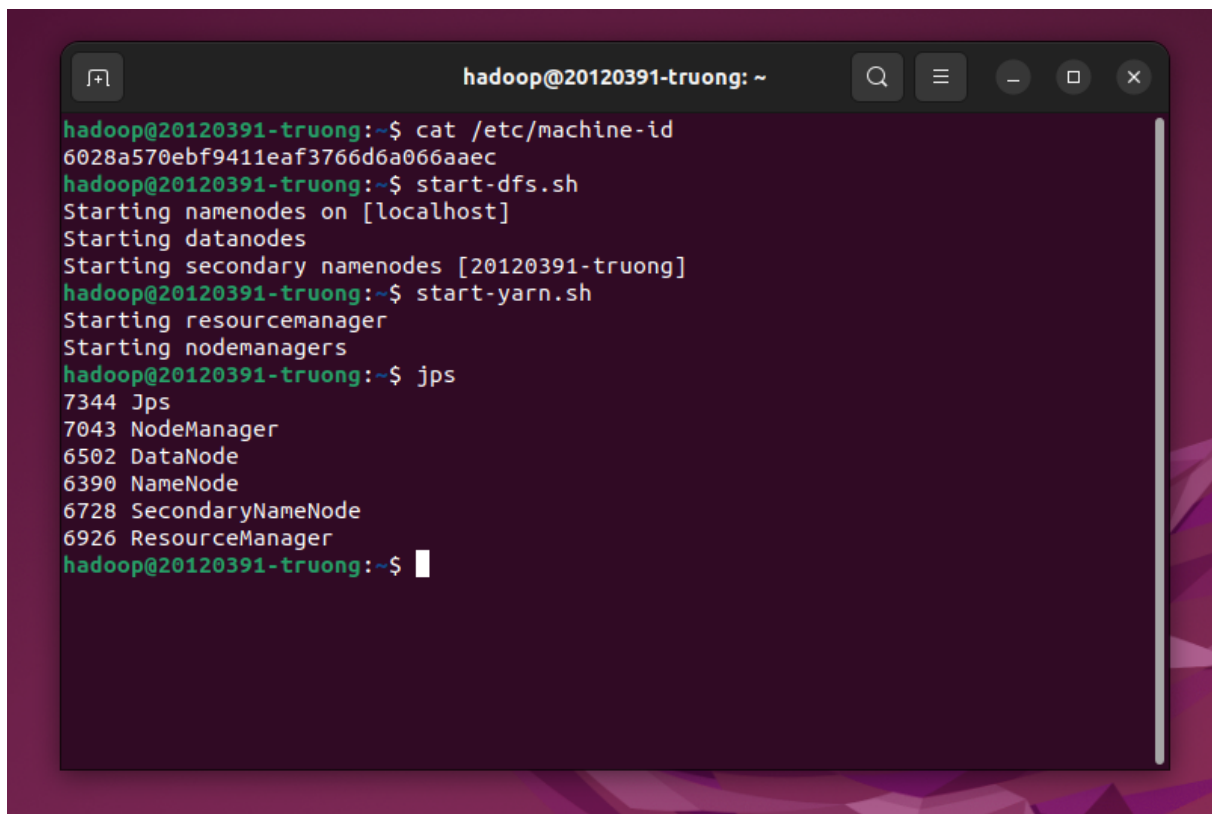
Figure 1.13: 20120366 done

1.5 Introduction to MapReduce

1. How do the input keys-values, the intermediate keys-values, and the output keys-values relate?

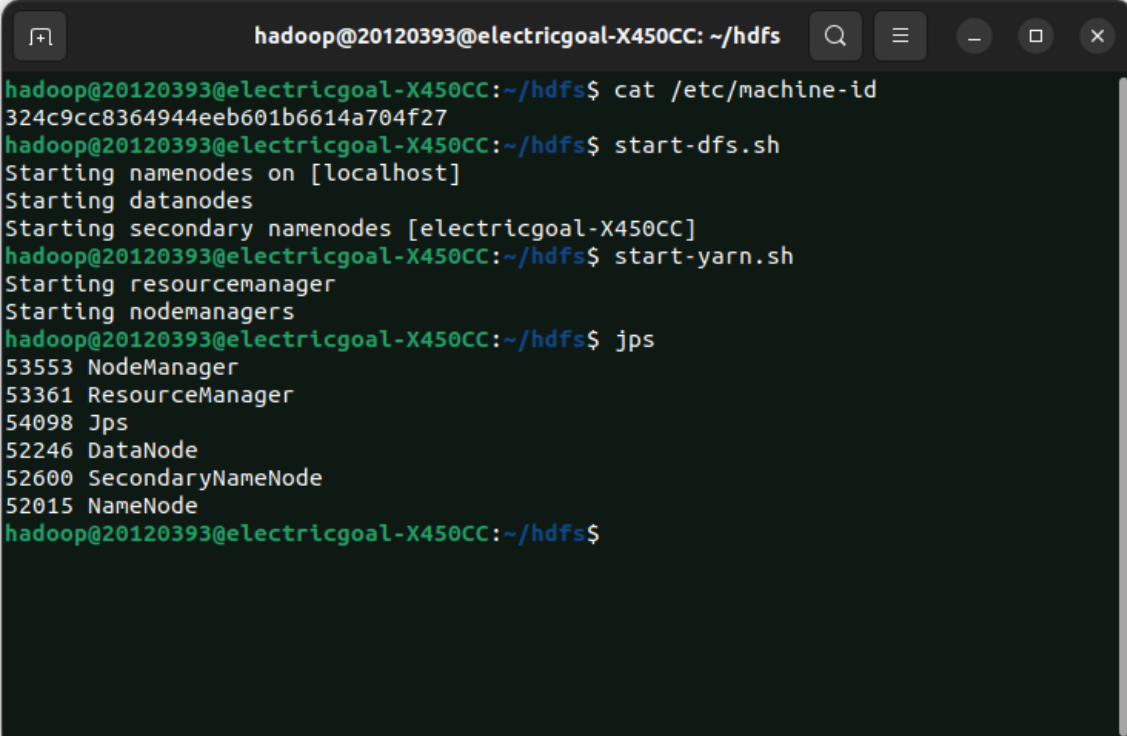
Answer :

- Input keys-values: The input data is divided into splits and represented as key-value pairs. Each input key-value is read into the MapReduce job using a RecordReader, which is responsible for reading the input data and converting it into key-value pairs.
- Intermediate keys-values: The map function processes the input keys-values and generates intermediate key-value pairs. The intermediate keys and values may be different from the input keys-values, depending on how the map function processes the data. The intermediate key-value pairs are sorted and grouped by key before being passed to the reduce function.
- Output keys-values: The reduce function generates the final output keys-values based on the intermediate key-value pairs that are passed to it. The output keys-values may be different from the intermediate keys-values, depending on how the reduce function processes the data. The output keys-values are typically written to a distributed file system, such as HDFS, or to a database. The output of the MapReduce job can be used as input to other MapReduce jobs or as input to other applications.

A terminal window titled 'hadoop@20120391-truong: ~' with standard window controls. The terminal shows the execution of several commands to set up Hadoop. First, 'cat /etc/machine-id' is run, displaying a long hexadecimal string. Then, 'start-dfs.sh' is executed, showing progress for starting namenodes on localhost, datanodes, and secondary namenodes on the current machine. Next, 'start-yarn.sh' is run, showing progress for starting the ResourceManager and NodeManagers. Finally, 'jps' is run to list the active Java processes, showing the Jps command itself and the NodeManager, DataNode, NameNode, SecondaryNameNode, and ResourceManager. The prompt returns to the shell after each command.

```
hadoop@20120391-truong:~$ cat /etc/machine-id
6028a570ebf9411eaf3766d6a066aaec
hadoop@20120391-truong:~$ start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [20120391-truong]
hadoop@20120391-truong:~$ start-yarn.sh
Starting ResourceManager
Starting nodemanagers
hadoop@20120391-truong:~$ jps
7344 Jps
7043 NodeManager
6502 DataNode
6390 NameNode
6728 SecondaryNameNode
6926 ResourceManager
hadoop@20120391-truong:~$
```

Figure 1.14: 20120391 done



```
hadoop@20120393@electricgoal-X450CC: ~/hdfs
hadoop@20120393@electricgoal-X450CC:~/hdfs$ cat /etc/machine-id
324c9cc8364944eeb601b6614a704f27
hadoop@20120393@electricgoal-X450CC:~/hdfs$ start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [electricgoal-X450CC]
hadoop@20120393@electricgoal-X450CC:~/hdfs$ start-yarn.sh
Starting resourcemanager
Starting nodemanagers
hadoop@20120393@electricgoal-X450CC:~/hdfs$ jps
53553 NodeManager
53361 ResourceManager
54098 Jps
52246 DataNode
52600 SecondaryNameNode
52015 NameNode
hadoop@20120393@electricgoal-X450CC:~/hdfs$
```

Figure 1.15: 20120393 done

2. How does MapReduce deal with node failures?

Answer :

Worker failure: The master node send heartbeat to each worker node. If a worker node fails, the master reschedule the tasks handled by the worker.

Master failure: The whole MapReduce job gets restarted through a different master based on check-pointed state of the failed master.

3. What is the meaning and implication of locality? What does it use?

Answer :

The concept of locality in the MapReduce refers to the idea that it is beneficial to process data on the same node where the data is stored, rather than moving it across the network to another node for processing. This is known as data locality.

MapReduce uses the concept of data locality to optimize the processing of data. The MapReduce framework is designed to distribute processing tasks to the nodes where the data is stored, in order to maximize data locality. When processing a large dataset, the framework splits the data into smaller chunks and distributes them across the cluster. Then, the Map tasks are scheduled on the same node where the data is stored, so that the data can be processed locally. Finally, the Reduce tasks are scheduled to aggregate the intermediate results generated by the Map tasks, again with the goal of minimizing data movement across the network.

4. Which problem is addressed by introducing a combiner function to the MapReduce model?

Answer :

The problem that is addressed by introducing a combiner function is the excessive duplicate data transfer during the shuffling phase of the MapReduce job. Without a combiner function, all the intermediate key-value pairs generated by the map tasks are transferred over the network to the reduce tasks, resulting in high network traffic and increased processing time.

By introducing a combiner function, the amount of data that needs to be transferred over the network is reduced, resulting in faster processing times and reduced network traffic. The combiner function helps to group together intermediate key-value pairs with the same key and perform a local aggregation, reducing the number of key-value pairs that need to be transferred. This is particularly useful when the same intermediate key appears multiple times across the map outputs.

1.6 Running a warm-up problem: Word Count

Use Eclipse IDE to run MapReduce on Ubuntu

1.6.1 Step 0: Install Eclipse on Ubuntu (if you had installed, please go to next step)

```
sudo snap install --classic eclipse
```

1.6.2 Step 1: Create new Java project

Open Eclipse, select **File** -> **New** -> **Java project**

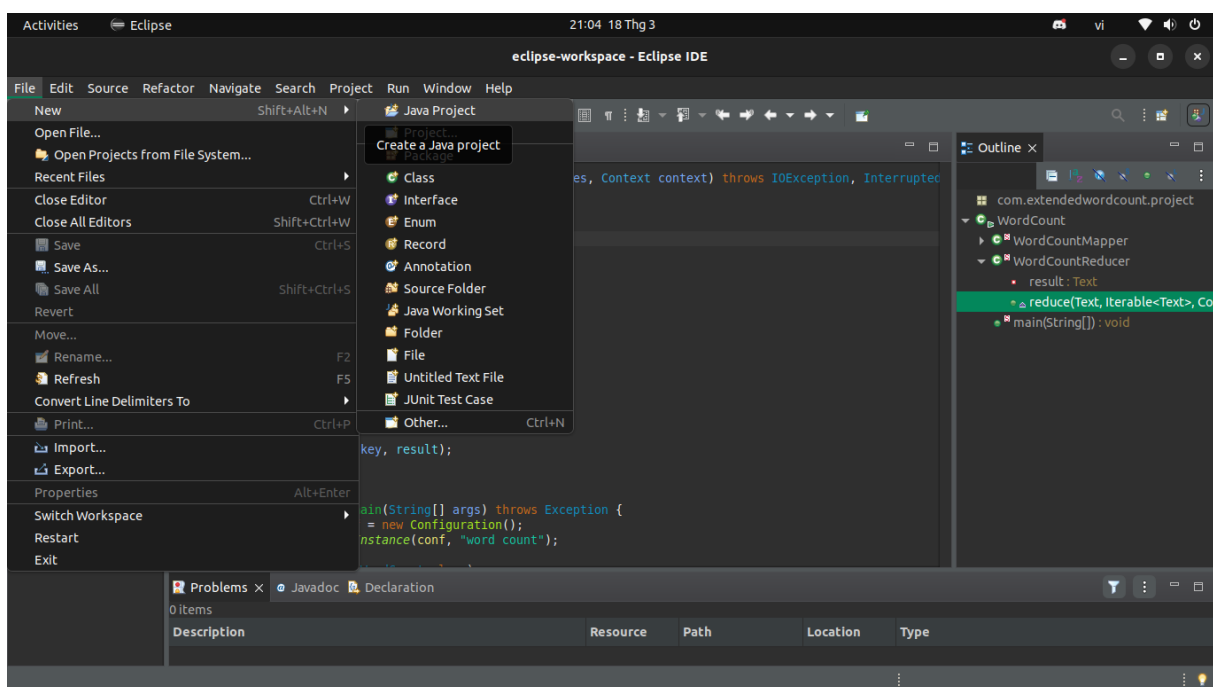


Figure 1.16: Run MapReduce

Enter project name and click on **Next** button

Click on **Finish** button

Result looks like this

1.6.3 Step 2: Delete file *module-info.java*

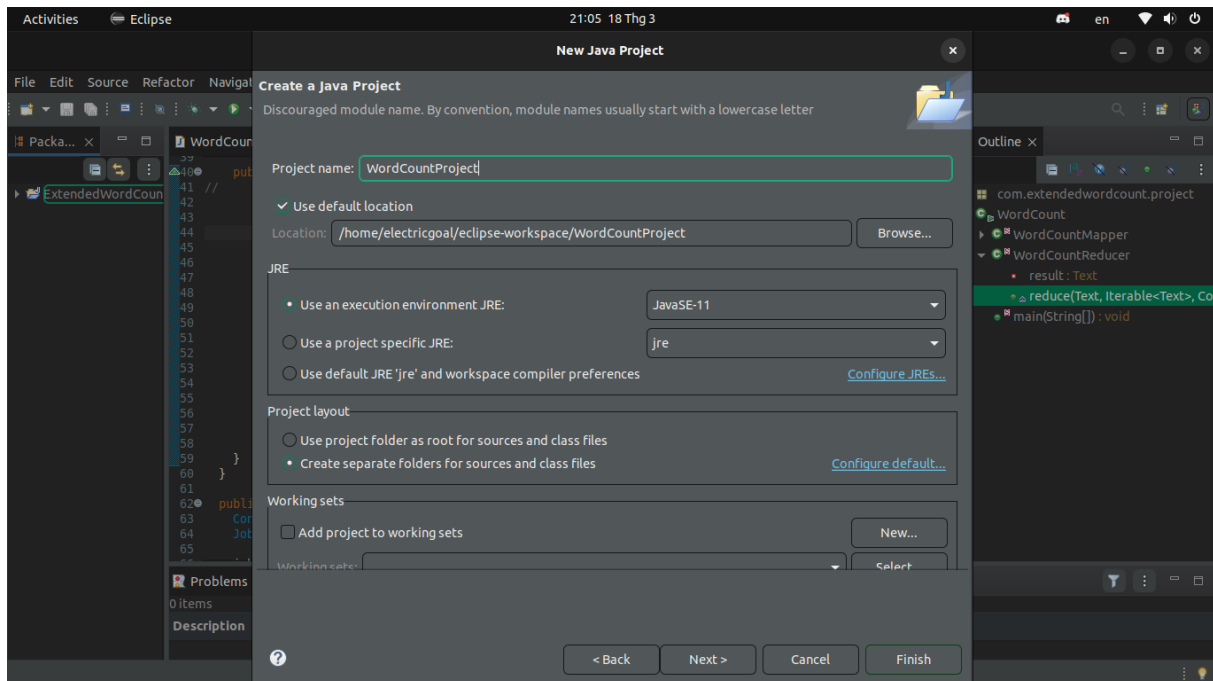


Figure 1.17: Run MapReduce

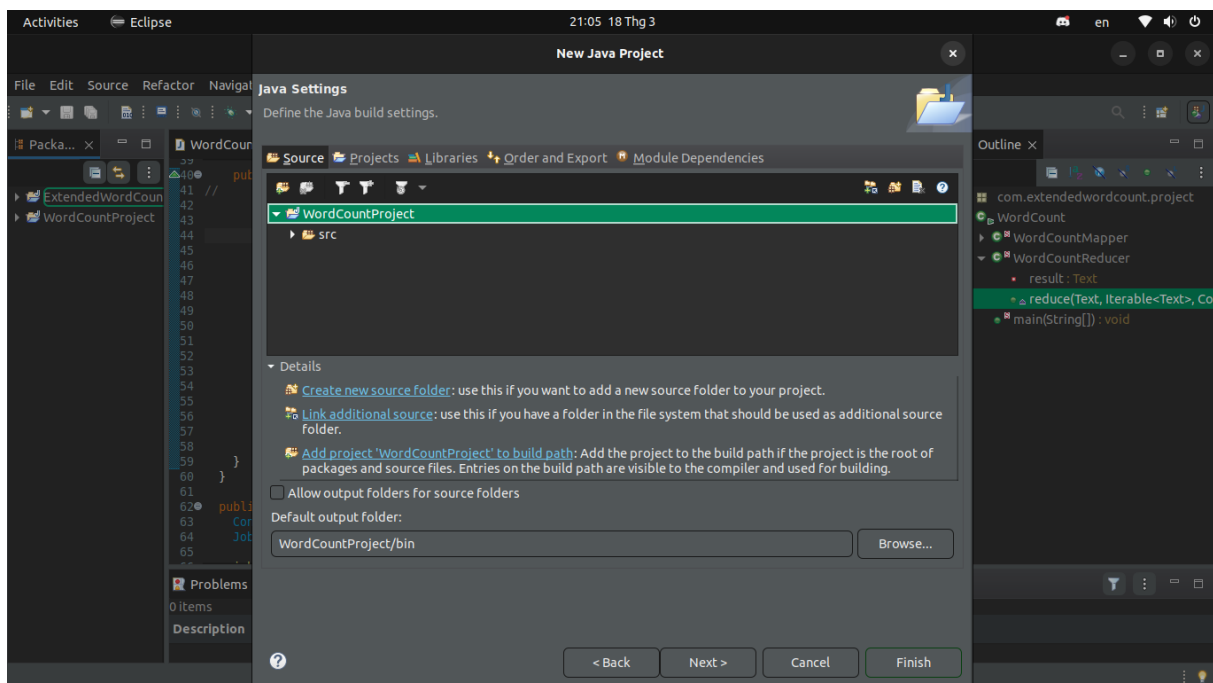




Figure 1.18: Run MapReduce

**Figure 1.19:** Run MapReduce**Figure 1.20:** Run MapReduce

1.6.4 Step 3: Create Java package

Right click on project name, select **New** -> **Package**

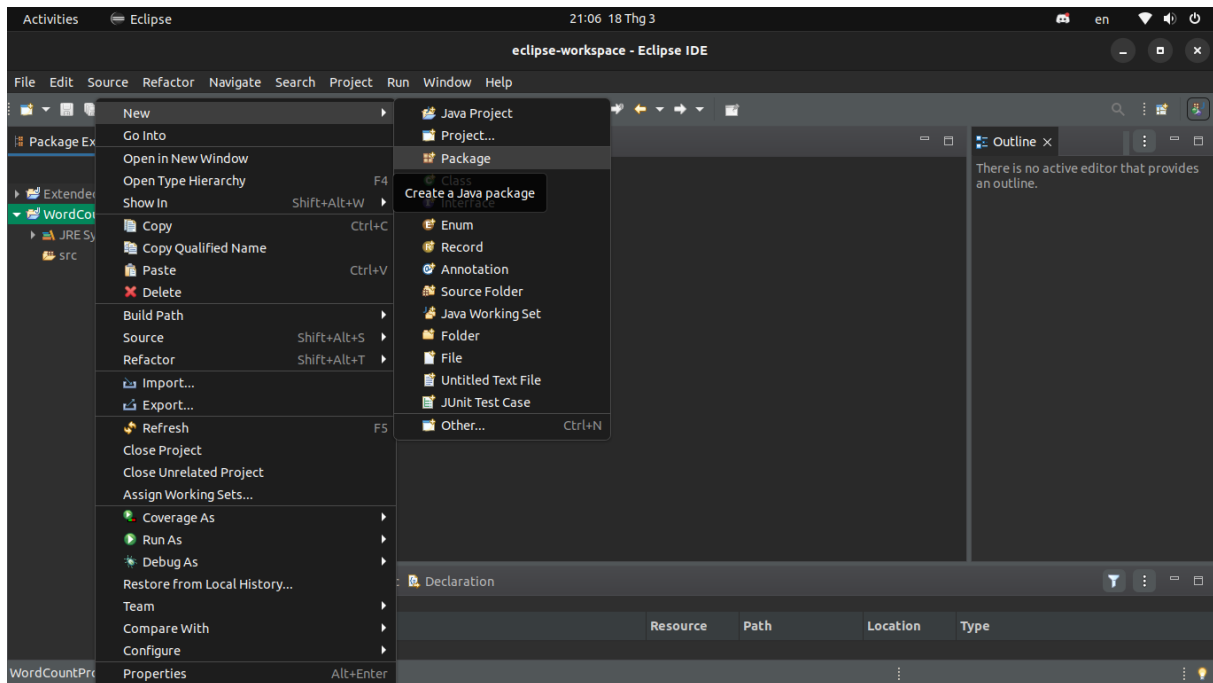


Figure 1.21: Run MapReduce

Enter Package name and click on **Finish** button

1.6.5 Step 4: Create Java class

Right click on project name, select **New** -> **Class** to create a Java class

Enter Class name and click on **Finish** button

1.6.6 Step 5: Paste WordCount code to the *WordCount.java* file just created

You should see many errors

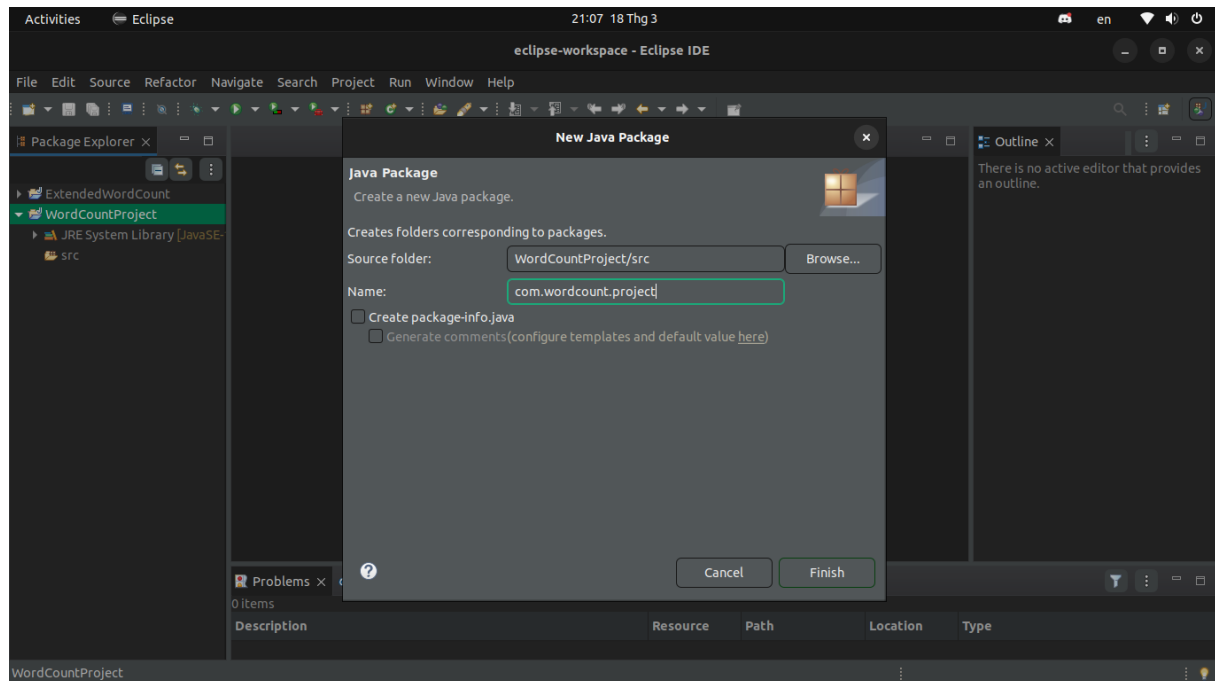


Figure 1.22: Run MapReduce

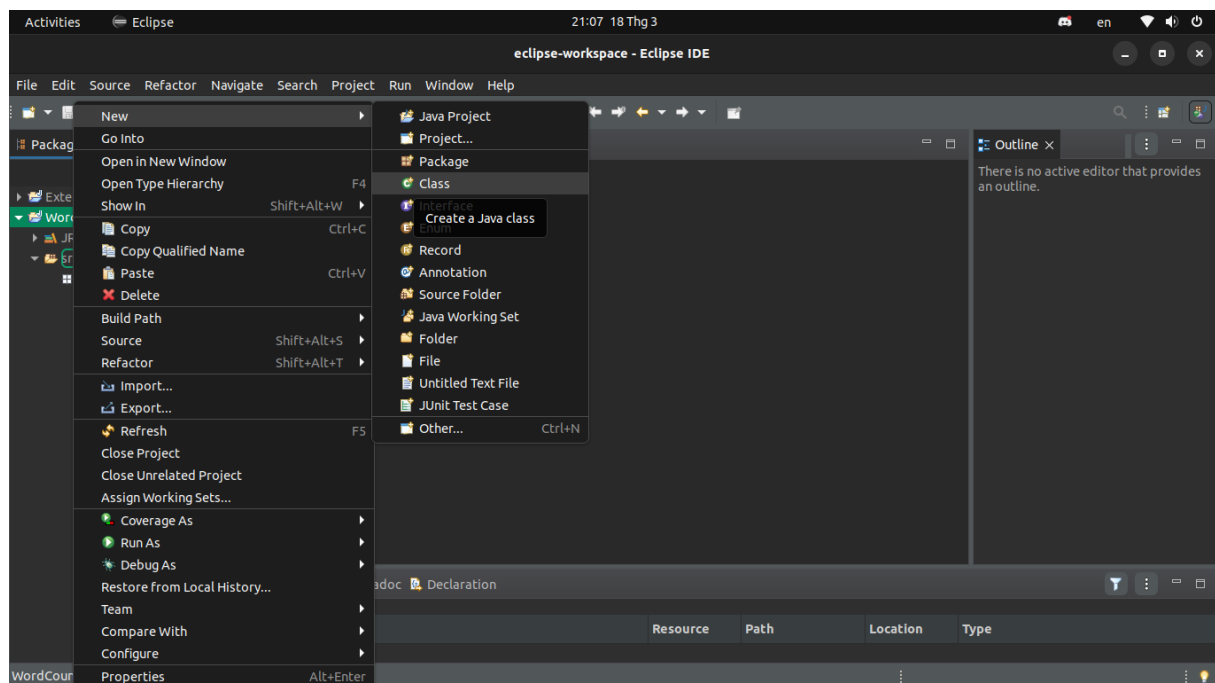


Figure 1.23: Run MapReduce

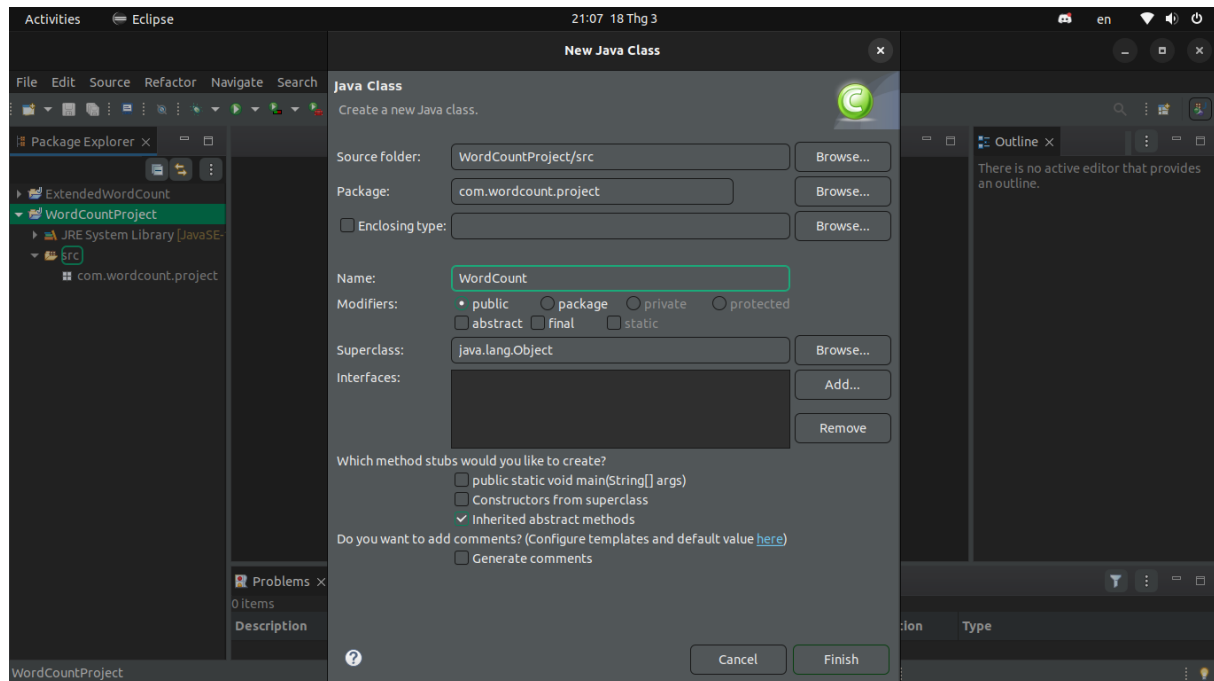


Figure 1.24: Run MapReduce

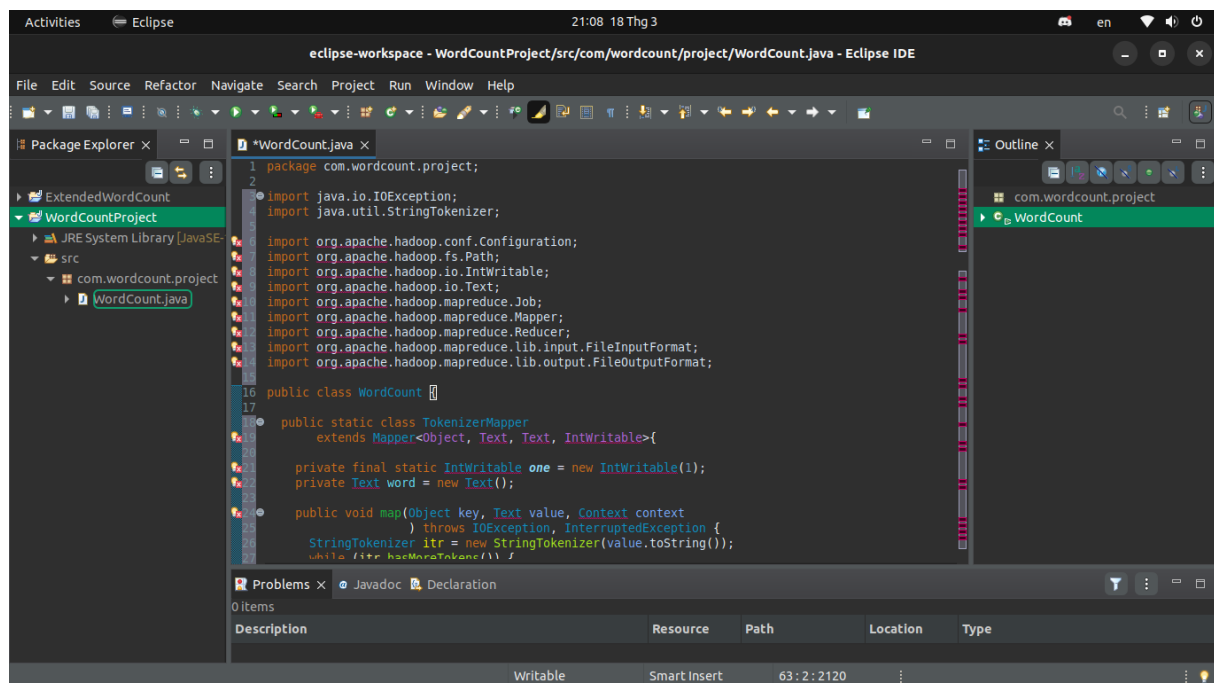


Figure 1.25: Run MapReduce

1.6.7 Step 6: Configure build path for the project

Right click on project name, select **New -> Build Path -> Configure Build Path**

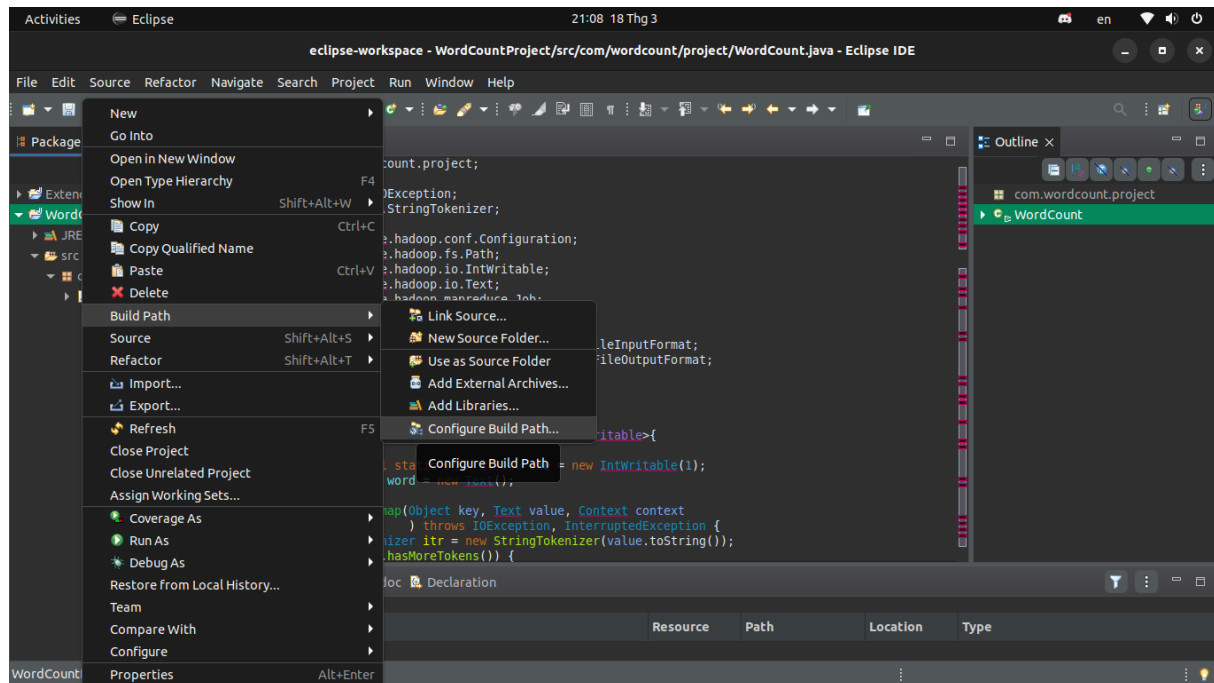


Figure 1.26: Run MapReduce

Click on the **Libraries** tab

Select **Classpath** section and click on the **Add External JARs** button

Navigate to the Hadoop installation directory and select the following JAR files:



- hadoop-mapreduce-client-core-<version>.jar
- hadoop-mapreduce-client-common-<version>.jar
- hadoop-mapreduce-client-jobclient-<version>.jar
- hadoop-common-<version>.jar

Click on the button **Apply and Close**

After that, the errors should disappear

1.6.8 Step 7: Export to JAR file

Right click to project name, select **Export**. You should see this screen, click on **JAR file -> Next**

**Figure 1.27:** Run MapReduce**Figure 1.28:** Run MapReduce




Figure 1.29: Run MapReduce

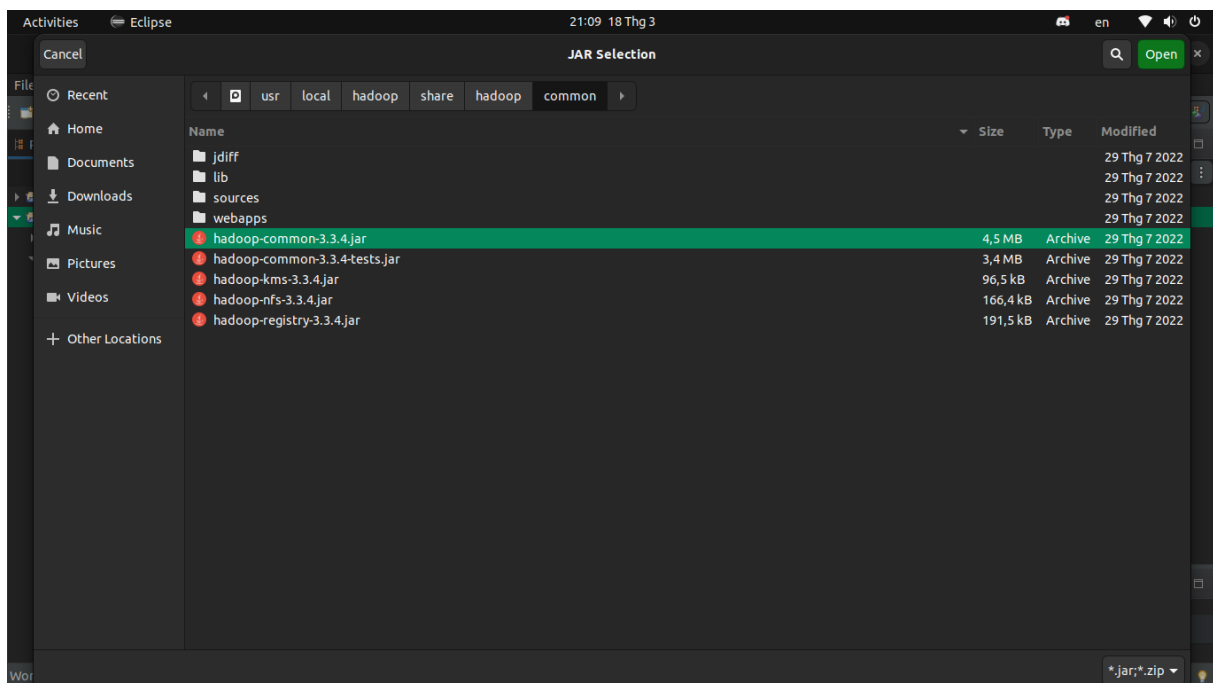


Figure 1.30: Run MapReduce

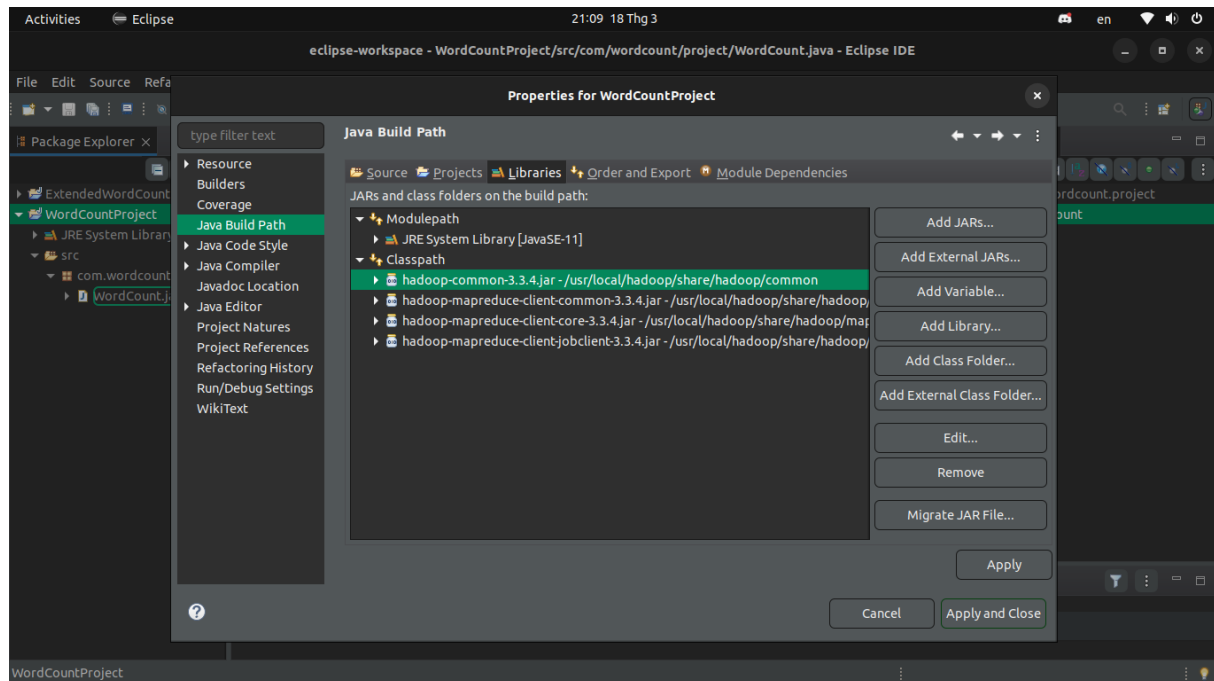


Figure 1.31: Run MapReduce

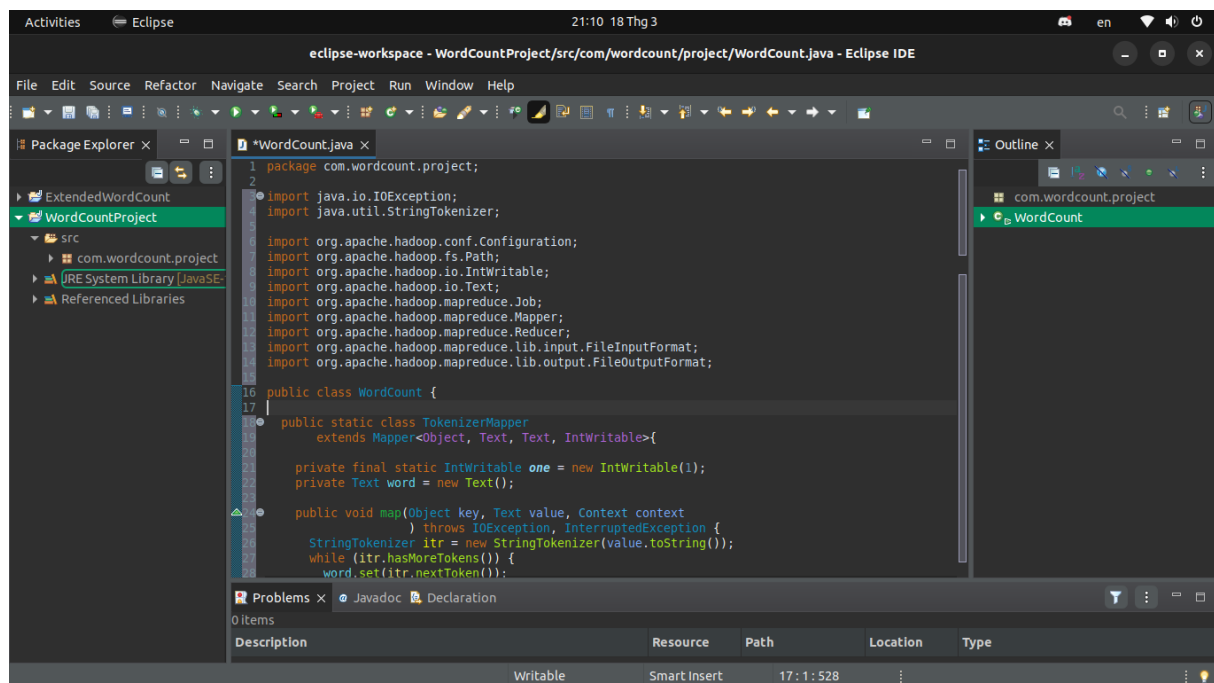


Figure 1.32: Run MapReduce

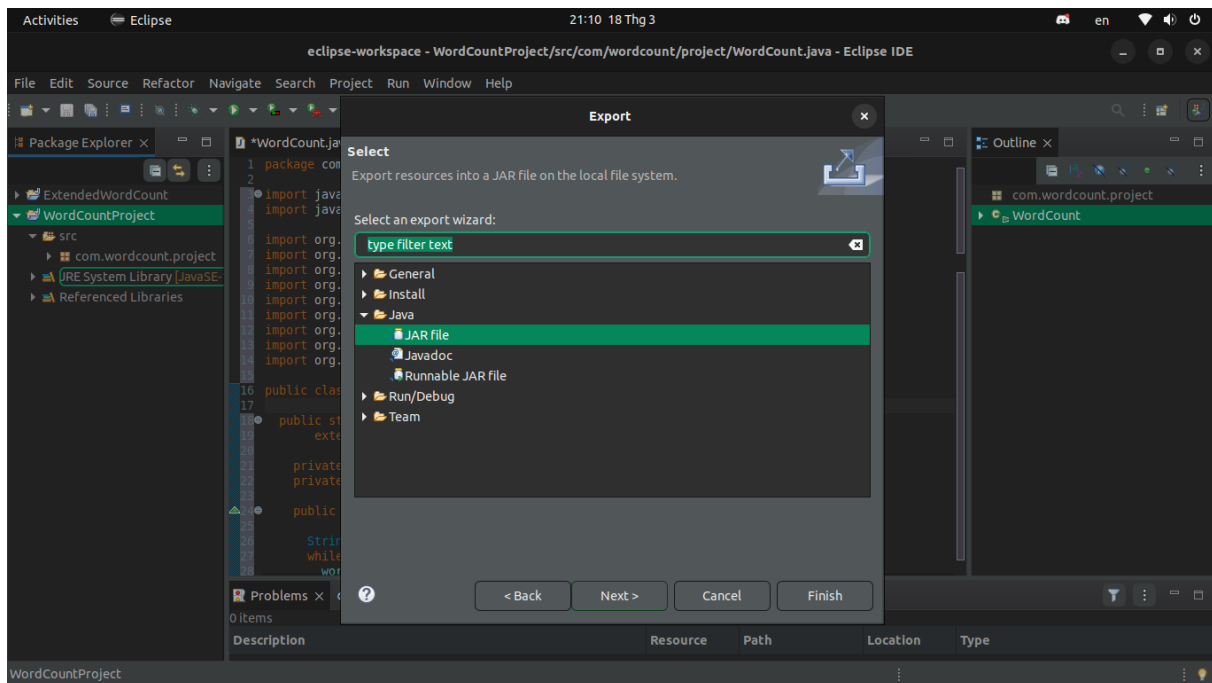


Figure 1.33: Run MapReduce

Enter name of jar file and path to save this jar file and. Once done, click on **Next** button

Click on **Next** button until see this screen and browse the the package in this project. Once done, click on **Finish** button

After all, you will get the Jar file

1.6.9 Step 8: Prepare to run MapReduce

Create new folder name “wordcount” in HDFS

```
hadoop fs -mkdir -p /<your-favorite-path>/wordcount
```

Create “input” folder in “wordcount” folder to store input file

```
hadoop fs -mkdir -p /<your-favorite-path>/wordcount/input
```

Put *input.txt* file into “input” directory

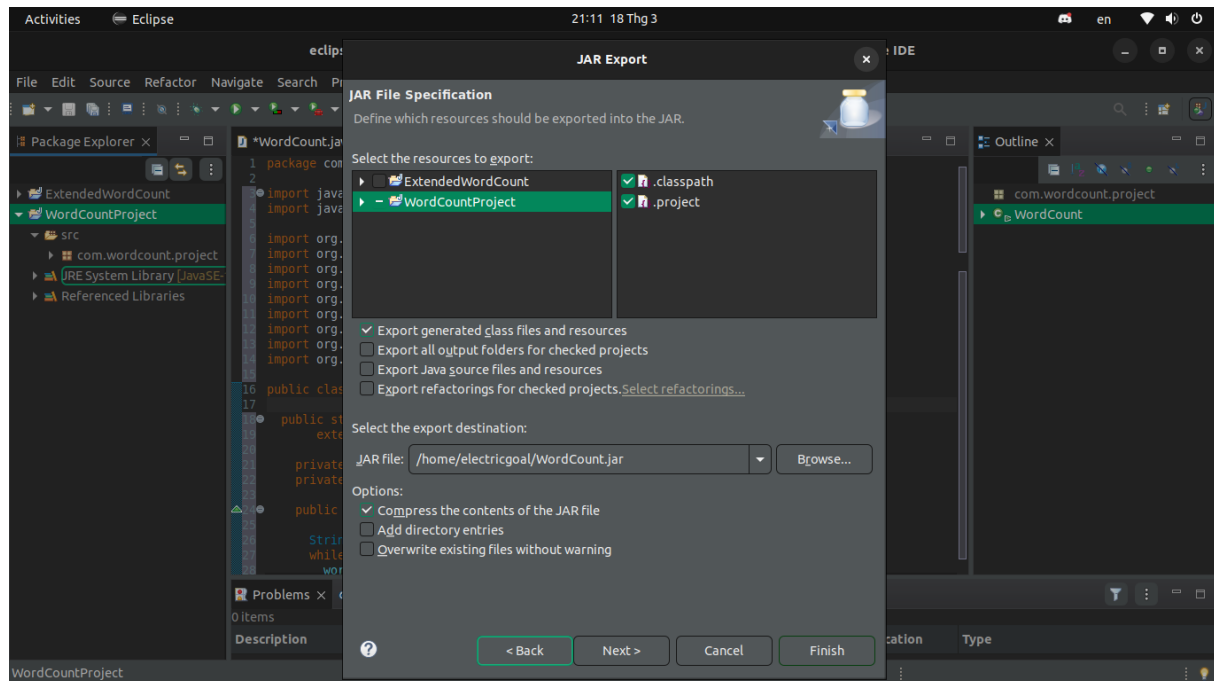


Figure 1.34: Run MapReduce

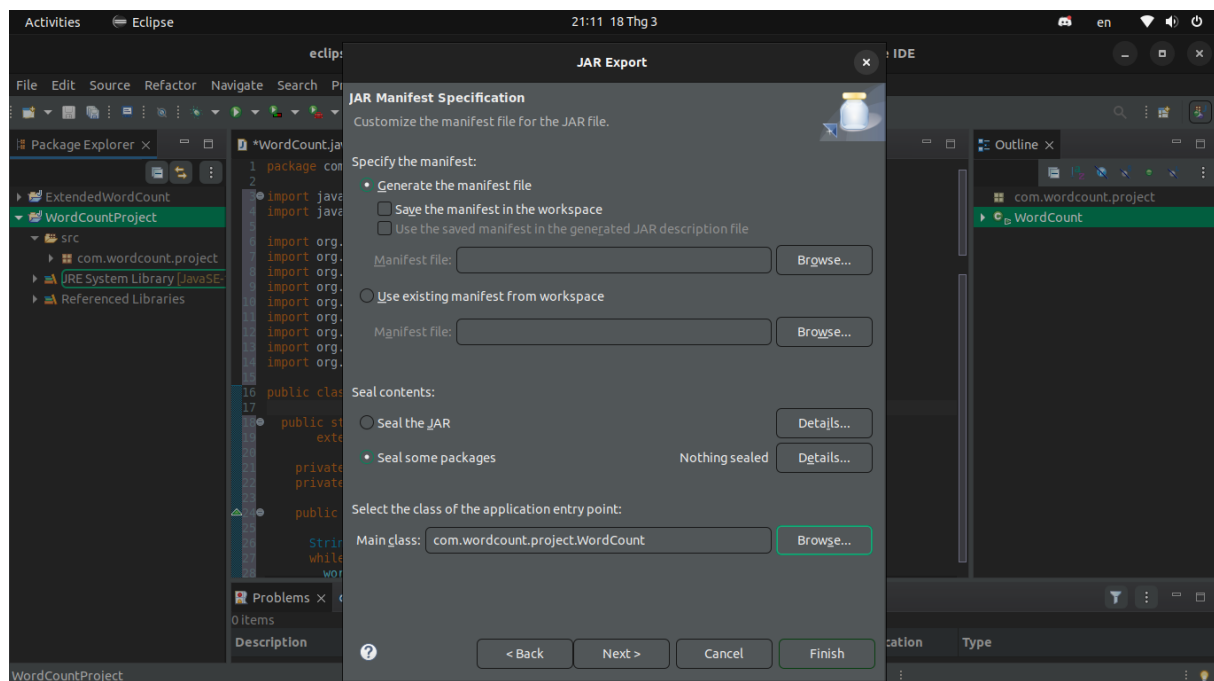


Figure 1.35: Run MapReduce

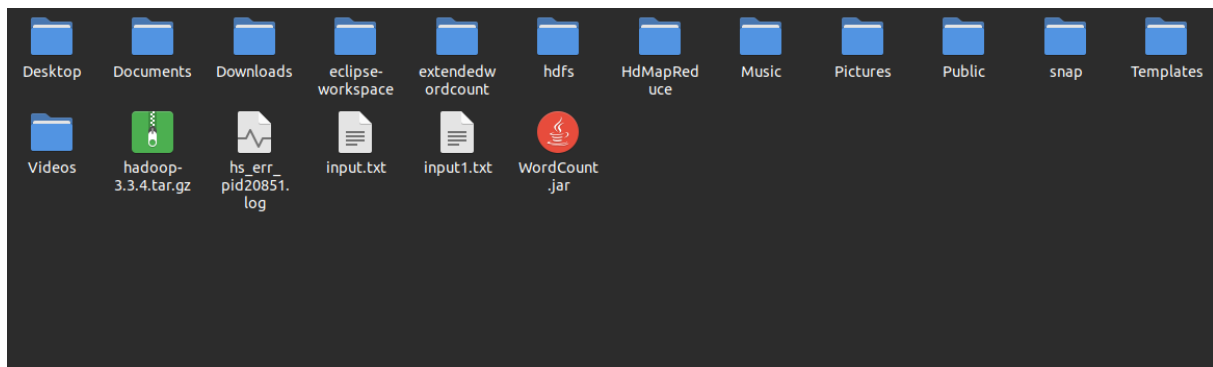



Figure 1.36: Run MapReduce

```
hadoop fs -put /<local_file_path>/input.txt /<your-favorite-path>/wordcount/input.txt
```

Open browser and enter `http://localhost:9870`, you should see the screen like this



Overview 'localhost:9000' (✓active)

Figure 1.37: Run MapReduce

Click on **Utilities** tab -> **Browse the file system**




Figure 1.38: Run MapReduce

Browse to your “wordcount” directory, you should see “input” folder. Click on it you will see **input.txt** file

Browse Directory

/user/hadoop/wordcount/input

Show 25 entries Search:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	hadoop	supergroup	172 B	Mar 15 22:03	1	128 MB	input.txt
-rw-r--r--	hadoop	supergroup	53 B	Mar 15 22:03	1	128 MB	input1.txt

Showing 1 to 2 of 2 entries

Previous 1 Next

Hadoop, 2022.

Figure 1.39: Run MapReduce

1.6.10 Step 9: Run MapReduce

```
hadoop jar WordCount.jar /<your-favorite-path>/wordcount/input/input.txt /<your-favorite-path>/wordcount/output
```

In my case, <your-favorite-path> is user/hadoop

```
hadoop jar WordCount.jar /user/hadoop/wordcount/input/input.txt /user/hadoop/wordcount/output
```

You should see something like this

To see the result, enter this command

```
hadoop fs -cat /<your-favorite-path>/wordcount/output/part-r-00000
```

In my case

```
hadoop fs -cat /user/hadoop/wordcount/output/part-r-00000
```

Compare to the input

1.7 Bonus

1.7.1 4.1 Extended Word Count: Unhealthy relationships

For more details, open folder src

```
Total megabyte-milliseconds taken by all reduce tasks=4298752
Map-Reduce Framework
  Map input records=22
  Map output records=22
  Map output bytes=238
  Map output materialized bytes=240
  Input split bytes=124
  Combine input records=22
  Combine output records=18
  Reduce input groups=18
  Reduce shuffle bytes=240
  Reduce input records=18
  Reduce output records=18
  Spilled Records=36
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=73
  CPU time spent (ms)=1580
  Physical memory (bytes) snapshot=500559072
  Virtual memory (bytes) snapshot=5532372992
  Total committed heap usage (bytes)=394264576
  Peak Map Physical memory (bytes)=293187584
  Peak Map Virtual memory (bytes)=2761269248
  Peak Reduce Physical memory (bytes)=207372288
  Peak Reduce Virtual memory (bytes)=2771103744
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=172
File Output Format Counters
  Bytes Written=162
hadoop@20120393@electricgoal-X450CC ->
```

Figure 1.40: Run MapReduce

```
hadoop@20120393@electricgoal-X450CC -> hadoop fs -cat /user/hadoop/wordcount/output/part-r-00000
Arianne 1
Belinda 1
Calantha 1
Christabel 1
Daria 1
Elizabeth 1
Evelyn 1
Helene 2
Lily 1
Maria 1
Mia 1
Nicole 1
Phoebe 1
Quinn 1
Susan 2
Udeline 1
Vivian 2
Zoe 2
hadoop@20120393@electricgoal-X450CC ->
```

```
1 Arianne
2 Calantha
3 Christabel
4 Daria
5 Udeline
6 Vivian
7 Elizabeth
8 Helene
9 Helene
10 Belinda
11 Lily
12 Vivian
13 Maria
14 Nicole
15 Phoebe
16 Susan
17 Zoe
18 Mia
19 Zoe
20 Quinn
21 Evelyn
22 Susan
```

Figure 1.41: Run MapReduce

Sample input:

```
faker showmaker  
gumayusi deft  
keria kellin  
canyon oner  
zeus canna  
chovy faker  
canyon peanut  
oner peanut  
zeus doran  
gumayusi peyz  
keria delight  
deft peyz  
delight kellin  
chovy showmaker  
doran canna
```

Expected putput:

```
canna neg  
canyon pos  
chovy pos  
deft eq  
delight eq  
doran eq  
faker eq  
gumayusi pos  
kellin neg  
keria pos  
oner eq  
peanut neg  
peyz neg  
showmaker neg  
zeus pos
```

The result screen:

```

hadoop@20120393@gotcha ~-> hadoop fs -cat unhealthy_relationship/output/part-r-00000
canna      neg
canyon     pos
chovy      pos
deft       eq
delight    eq
doran      eq
faker      eq
gumayusi   pos
kellin     neg
keria      pos
oner       eq
peanut     neg
peyz       neg
showmaker  neg
zeus       pos

```

Figure 1.42: Result

1.7.2 4.2 Setting up Fully Distributed Mode

Uncomplete

1.8 References

- Example: WordCount v1.0: https://hadoop.apache.org/docs/current/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html#Example:_WordCount_v1.0
- How to Install Apache Hadoop on Ubuntu 22.04: <https://www.howtoforge.com/how-to-install-apache-hadoop-on-ubuntu-22-04/>
- How to run Word Count example on Hadoop MapReduce (WordCount Tutorial): <https://www.youtube.com/watch?>
- MapReduce Word Count Example using Hadoop and Java: <https://www.youtube.com/watch?v=qgBu8Go1SyM>
- Create and Execute your First Hadoop MapReduce Project in Eclipse: <https://medium.com/data-science-community-srm/create-execute-your-first-hadoop-mapreduce-project-with-eclipse-9ec03105e974>
- All of StackOverflow link related:
 - <https://stackoverflow.com/questions/11889261/datanode-process-not-running-in-hadoop>
 - <https://stackoverflow.com/questions/66182686/why-output-key-value-of-mapper-needs-to-be-same-as-that-of-output-key-value-ofco>